



EagleEye

Large photoset visualization

Carlos Eduardo Henriques de Jesus Fonseca

Dissertação para a obtenção de Grau de Mestre em
Engenharia Informática e de Computadores

Júri

Presidente: Nome do Presidente
Orientador: Nome do Orientador
Co-orientador: Nome do Co-orientador
Vogais:
Nome do Vogal 1
Nome do Vogal 2
Nome do Vogal 3

Mês e Ano

Dedicated to someone special...

Acknowledgments

A few words about the university, financial support, research advisor, dissertation readers, faculty or other professors, lab mates, other friends and family...

Resumo e palavras chave

Inserir o resumo em Português aqui com o máximo de 250 palavras e acompanhado de 4 a 6 palavras-chave...

cenas

Palavras-chave: palavra-chave1, palavra-chave2,...

Abstract and keywords

Insert your abstract here with a maximum of 250 words, followed by 4 to 6 keywords...

Keywords: keyword1, keyword2,...

Contents

Acknowledgments	v
Resumo e palavras chave	vii
Abstract and keywords	ix
List of Tables	xiii
List of Figures	xv
Acronyms	xvii
1 Introduction	1
1.1 Prologue	1
1.2 Context	1
1.3 Motivation	1
2 Related Work	2
2.1 Visual guided navigation for image retrieval	2
2.2 Does organisation by similarity assist image browsing?	3
2.3 Browsing large collections of images through unconventional visualization techniques	4
2.4 A comparison of static and moving presentation modes for image collections	4
2.5 Organizing and browsing photos using different feature vectors and their evaluations	5
2.6 An evaluation of colour-spatial retrieval techniques for large image databases	6
2.7 Automatic organization for digital photographs with geographic coordinates	6
2.8 Similarity pyramids for browsing and organisation of large image databases	7
2.9 NN ^k networks for content-based image retrieval	9
2.10 Phorigami: A Photo browser based on meta-categorization and origami visualization	10
2.11 A next generation browsing environment for large image repositories	11
2.12 Flexible access to photo libraries via time place, tags, and visual features	12
2.13 Discussion	13
3 Solution Requirements	15
3.1 Motivation	15
3.2 Goals	15
3.2.1 Main Goal	16
3.2.2 Design Goal	16

4 Eagle Eye	17
4.1 Architecture	17
4.1.1 Design Decisions	17
4.1.2 Visualization	18
4.1.3 Backend	18
4.2 Backend	18
4.2.1 Structure	18
4.3 Interface	18
5 Evaluation	19
6 Conclusions	21
6.1 Achievements	21
6.2 Future Work	21
Bibliography	25

List of Tables

2.1 Different browsing methods	14
--	----

List of Figures

2.1	The chromacy diagram is split in parts and each image belongs to one of this parts. The diagram is part of the user interface (UI) and when navigating through the diagram, only the images related to that part are shown.	2
2.2	Three arrangements of 100 images of Kenya, based on visual similarity. On the left is the arrangement with overlap, in the middle a 12x12 grid (which removes the overlap while preserving some of the structure), and on the right a 10x10 grid (which maximises the thumbnail size).	3
2.3	Spot display.	4
2.4	Shot display.	4
2.5	The six rapid serial visual presentation modes used in the experiments	5
2.6	The result obtained for organizing 2200+ photos using colour autocorrelogram feature vector, using [36]	6
2.7	Images organised with [9]. Images with similar colour and texture are spatially adjacent. .	7
2.8	Local network around the chosen butterfly image depicted in the middle.	9
2.9	On the left, an example of an interaction on a group of photos that makes a panorama. On the right, a visualisation on 537 photos with some groups.	10
2.10	Hue sphere of a dataset	11
2.11	Photos Grouped by Geographic Similarity and Filtered by Date and Place.	12

Acronyms

UI user interface

SOM self organising map

CBIR content based image retrieval

MP megapixel

FEP feature extraction plugin

WPF Windows Presentation Foundation

Chapter 1

Introduction

1.1 Prologue

Describe the past. Describe the main goal of the work Describe that this will show that the work is relevant

1.2 Context

Present concepts...?

1.3 Motivation

??

Chapter 2

Related Work

There has been a lot of discussion about the organisation and retrieval methods of large image libraries and some interesting and relevant are summarised here.

2.1 Visual guided navigation for image retrieval

Qiu et al. explore in [28] the requirements of a system intended for visualising large photo collections. They identify as the two most important requirements, the first being an easy to use UI, that gives clean information to the user and helps to create a mental image of the whole collection helping him to navigate on the collection. The second requirement is responsiveness because while image processing can be an heavy task, the user needs to be able to interact with the interface and he won't use the application if it's slow.

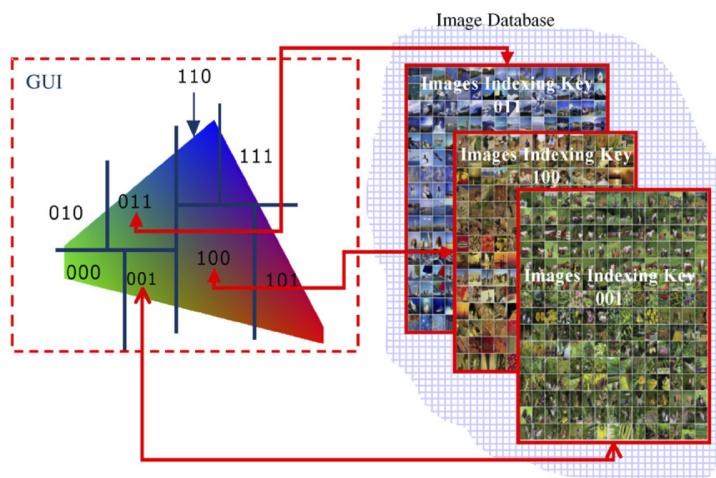


Figure 2.1: The chromacy diagram is split in parts and each image belongs to one of this parts. The diagram is part of the UI and when navigating through the diagram, only the images related to that part are shown.

The system shows all the photos arranged by colour, just as many others like it. The difference is the process in use. Instead of calculating distance vectors based on the histogram of each image, this

approach classifies each image with a simple description, like a mean of its colours, and arranges them by that value, on a colour map (fig. 2.1). The process is much faster but is also more error prone, specially on photos without a clear main colour.

Their tests show they achieved good responsiveness and better results than using a file explorer.

2.2 Does organisation by similarity assist image browsing?



Figure 2.2: Three arrangements of 100 images of Kenya, based on visual similarity. On the left is the arrangement with overlap, in the middle a 12x12 grid (which removes the overlap while preserving some of the structure), and on the right a 10x10 grid (which maximises the thumbnail size).

The aim of [31] by Rodden and Sinclair was to evaluate how photo organisation by similarity (fig. 2.2) could benefit a user looking for images. Some users tested an application that could show the same images both in a random and in an organised by similarity way. This organisation by similarity was based on a rough description of the images but it could be other descriptors.

The results differ if the user knows what he's looking for or not. In case he does, being able to filter only the relevant images makes it quick to find the ones that matter. This obviously depends on the quality of the labelling. Users reported that sometimes the similar images appear to merge.

In case the user doesn't know what he's looking for, the random approach might be helpful because the strong images usually contrast to their neighbours and thus appear to stand out.

For some people, having access to different arrangements of the same set of images is useful, although the source of the individual differences still needs to be determined.

2.3 Browsing large collections of images through unconventional visualization techniques

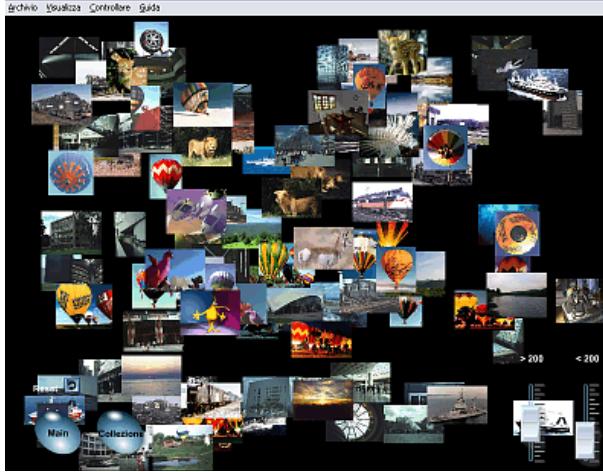


Figure 2.3: Spot display.



Figure 2.4: Shot display.

Porta describes in [27] a few visualisation methods for large collections of images and tests them with users. The purpose it to find ways or metaphors that provide a good visualisation experience in terms of time spent and quality of the visualisation.

Some of the various techniques were the simple image grid view, a grid view with variable and independent height and width (EIB), a view that animates images like they were shot from a distance and get closer to the user called Shot (fig. 2.4), a view where images quickly appear on random positions on screen named Spot (fig. 2.3), and some other less commons like one that simulates an cylinder created with images (Cylinder), and others less relevant (Rotor, Tornado and Tornado of Planes).

The testing was based on the efficiency of users searching for specific images on a collection of a thousand photos. The Spot view was the best, followed the Shot, Cylinder finally the common grid view. The other views got scores near or below the grid view.

2.4 A comparison of static and moving presentation modes for image collections

This paper [12] by Cooper et al. is not about large libraries but about what kind of interfaces for image showing has greater success of user recognition and preference (fig. 2.5).

With the help of eye-tracking techniques and user preference, the authors determined that static images are better than moving ones because makes them easier to recognise and avoid some user confusion.

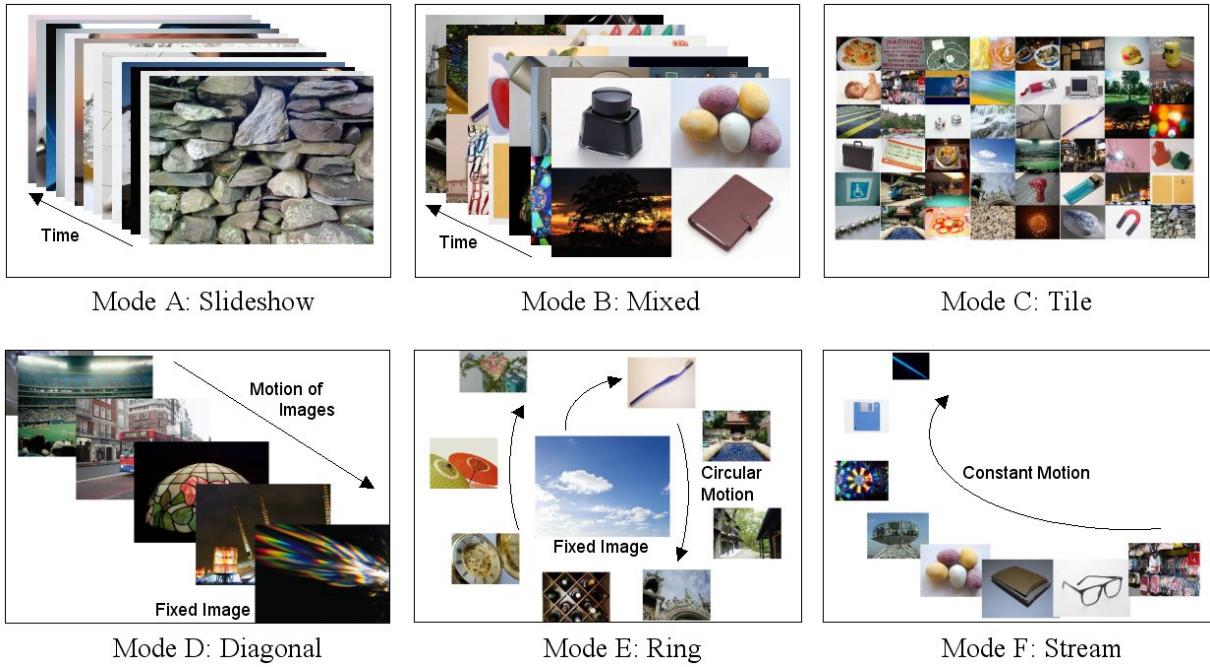


Figure 2.5: The six rapid serial visual presentation modes used in the experiments

2.5 Organizing and browsing photos using different feature vectors and their evaluations

Although it doesn't mention why, Strong [36] focuses on the better experience provided by colour organisation of a large image collection (fig. 2.6).

A self organising map (SOM) is used to display the images on the screen featuring zooming, panning and sorting capabilities. The work is then based around the various methods used to determine the images' similarity.

Simpler methods are based on colour histograms, which aren't affected by rotations or scales but, by not having spacial information on colours, allows very different images be closer together.

Other methods rely on gradients which contain spatial information and, therefore, are sensitive to image contents but not colour. In general, the best methods were found to be the hybrid ones, where both colour histograms and gradients are used to classify the images. No user testing is made in this project, neither is the speed of image categorisation referred about the methods used.



Figure 2.6: The result obtained for organizing 2200+ photos using colour autocorrelogram feature vector, using [36]

2.6 An evaluation of colour-spatial retrieval techniques for large image databases

Tan et al. [38] present an evaluation of three colour-spatial image retrieval techniques.

The signature-based technique creates a signature for each image, based on the most frequent colours, according to a threshold, of each subdivision, or bin, of that image. The comparison between images is made by comparing the main colours present on each bin. It is possible to assign more weight to specific bins according to the user's interest.

The partition-based approach is also based on bins, each having its own colour histogram. The similarity between images is given by the distance of the histograms of the corresponding bins.

The cluster-based method bases on the fact that humans focus on large patches (clusters) of the same colour and, therefore, two images will appear similar if both have similar coloured clusters on at roughly the same location. This method extracts the larger clusters and their colour from each image. The similarity is calculated by the amount of overlap between clusters.

This techniques were tested with a collection of 12,000 images and, besides the colour information, the brightness was also analysed for increased performance. The authors conclude the signature method was generally better on both effectiveness and efficiency.

2.7 Automatic organization for digital photographs with geographic coordinates

In this paper, by Naaman et al. [25], is described a system that organises digital photographies accordingly to location and date embedded on the metadata.

The objective was trying to mimic the way people think about their collections. Photos are usually bursts separated by some time. Based on this and on the different places, events can be created

to agglomerate photos from the same bursts. Location naming is done by calculating the most relevant places, like parks or cities, and then mixing the more precise locations with the more important neighbour cities to create a relevant and identifiable name. This was specially important since this work didn't involve showing any maps but only the location names and events.

The authors showed good results and, nowadays, some common applications use similar features although including maps.

2.8 Similarity pyramids for browsing and organisation of large image databases

Chen et al. present in [9] a method for designing a hierarchical browsing environment called a similarity pyramid. The similarity pyramid groups similar images together while allowing users to view the database at varying levels of resolution. Each level is organised such that similar images are in close proximity on a two-dimensional grid (fig. 2.7). Images are first organised into a binary tree through agglomerative clustering based on colour, edge and texture similarities. The binary tree is transformed into a quadtree, a tree in which each node has four children instead of only two.

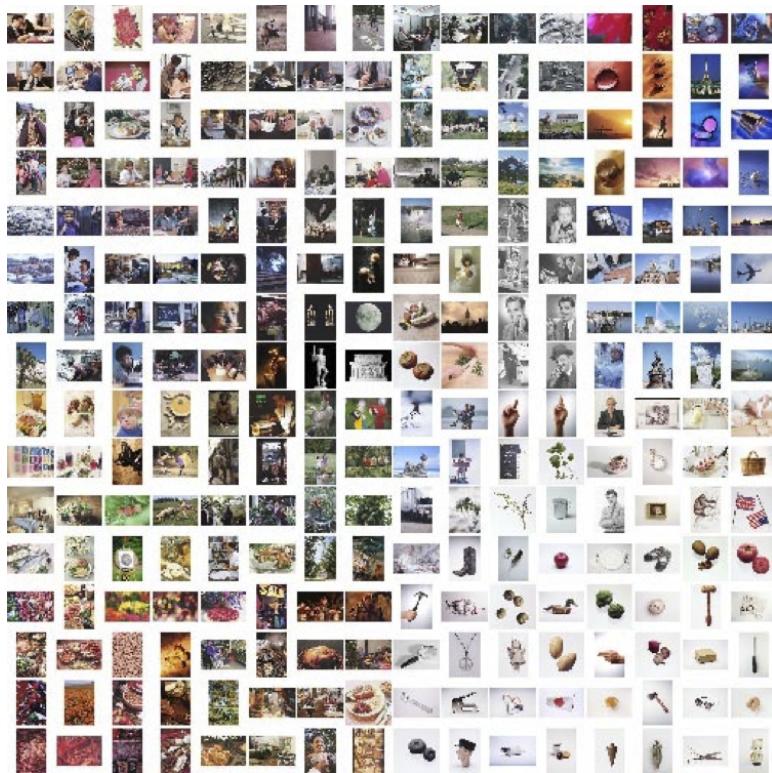


Figure 2.7: Images organised with [9]. Images with similar colour and texture are spatially adjacent.

The similarity pyramid is best constructed using agglomerative (bottom-up) clustering methods, and a fast-sparse clustering method is presented which dramatically reduces both memory and computation over conventional methods. This method is based on the flexible agglomerative clustering algorithm, but using only a sparse proximity matrix and exploiting the author's approximate branch and bound search

algorithm.

The authors found that the method for mapping the clustering to a pyramid can make a substantial difference in the quality of organisation. Finally, a dispersion metric for objectively measuring pyramid organisation was proposed, and found that it correlated well with the author's subjective evaluations of pyramid organisation.

2.9 NN^k networks for content-based image retrieval

Heesch describes in [18] a different interaction technique for content based search in large image collections. Each image is a vertex in a graph and arcs are established between images if there exists at least one combination of features for which one image is retrieved as the nearest neighbour of the other. Each arc is weighted by the proportion of feature combinations for which the nearest neighbour relationship holds. By integrating the retrieval results over several feature combinations, the resulting network helps expose the semantic richness of images.

The interface reflects the vertexes and respective arcs, allowing to browse between the related images (fig. 2.8) in the network.

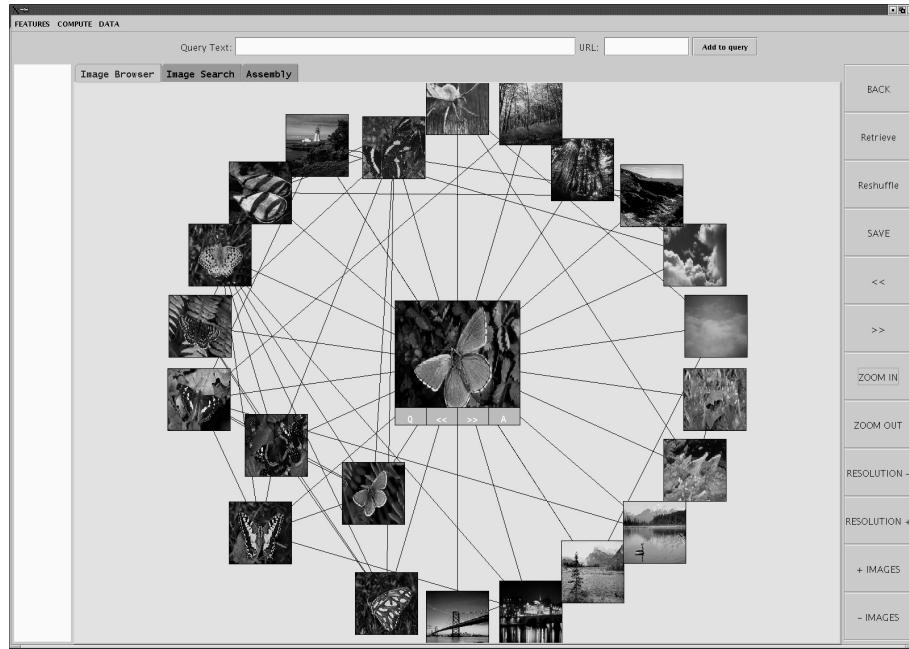


Figure 2.8: Local network around the chosen butterfly image depicted in the middle.

Seven low-level features are used for the classification of the images. HSV Global Colour Histograms maps the images by colour, saturation and brightness; Colour Structure Descriptor maps the distribution of colours by dividing each image in 64 windows, the colour space in 184 bins and associating colour bins with image windows; Thumbnail feature compares identical images by saving the grey value of each pixel from a scaled down version of the image; Convolution filters discovers very selective features by reapplying 25 filters three times; Variance feature calculates standard deviations within a sliding window; Uniformity Feature is another statistical feature, calculating the grey level of an image split in 64 parts; Bag of words is the last feature and weights words associated to each image.

Tests showed great results using a mix of search, relevance feedback and browsing, and even only browsing was considerably better than other, more restrictive, approaches.

The technique helps avoid the problem of image polysemy by showing all gathered meanings of the images to the user. The feature network is pre-computed, allowing for quick realtime browsing. The authors claim it took 50 hours to process 32000 images but make no reference to the possibility adding

images to the collection, after the computation.

2.10 Phorigami: A Photo browser based on meta-categorization and origami visualization

Hsu et al. [20] try to ease the browsing problem by analysing the collections and identifying groups of related pictures. Each type of group is visualised in a specific way, inspired by the Origami art.

Groups of similar or related photos were manually classified based on camera movement and subject movement, creating different types of groups static view where both camera and subject are fixed and is presented as a panorama; multi-view where the subject is fixed but the camera is moving and is shown as a presentation; if the subject is moving, the photos are categorised as motion capture and can be shown as an animated photo (fixed camera) or a presentation (moving camera); finally group photos, where different groups of people are photographed, are shown as a folding presentation.

This covers various cases where the photographer takes a few similar photos of the same subject because it's either a panorama, various angles or just to be sure the photo was well captured.

The interface implements the different presentation types as different metaphors, easy for the user to understand, like a folded paper on a wide panorama that can be expanded (fig. 2.9). Although some of them appear to be a little hard to distinguish in its compressed form, it shouldn't be difficult to make it clearer. Other possible problem is the use of different touch interactions for each presentation type that might confuse users on what gesture should they use.

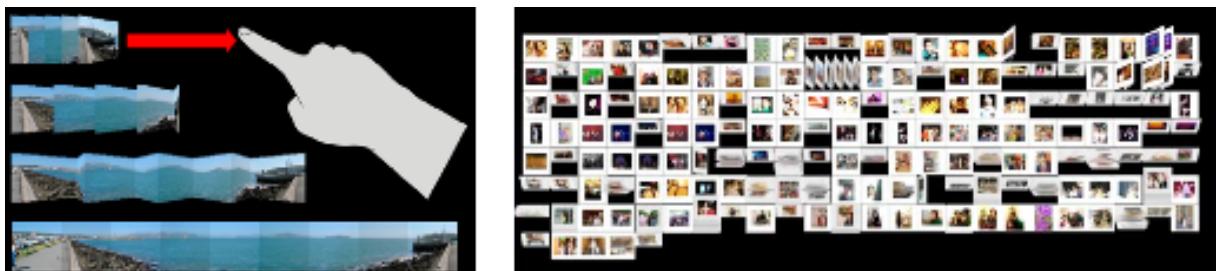


Figure 2.9: On the left, an example of an interaction on a group of photos that makes a panorama. On the right, a visualisation on 537 photos with some groups.

2.11 A next generation browsing environment for large image repositories



Figure 2.10: Hue sphere of a dataset

Schaefer [34] tries to take similarity based organisation of images from the 2D space to a 3D sphere, which allows interaction from the users. Rotating the sphere reveals images with different colours while tilting it reveals brighter or darker images.

Large image collections are handled through a hierarchical approach that brings up similar, previously hidden, images when zooming in on an area.

The description of the colour is retrieved by calculating the image's median colour for its efficiency over other methods like histograms. This two features are directly mapped onto the sphere's coordinates and the entire structure is pre-calculated so browsing can be performed in real-time. Image overlapping is also avoided (fig. 2.10).

The work was tested on a 4500 image collection with no evaluation as to its performance and a weak and subjective user testing.

2.12 Flexible access to photo libraries via time, place, tags, and visual features

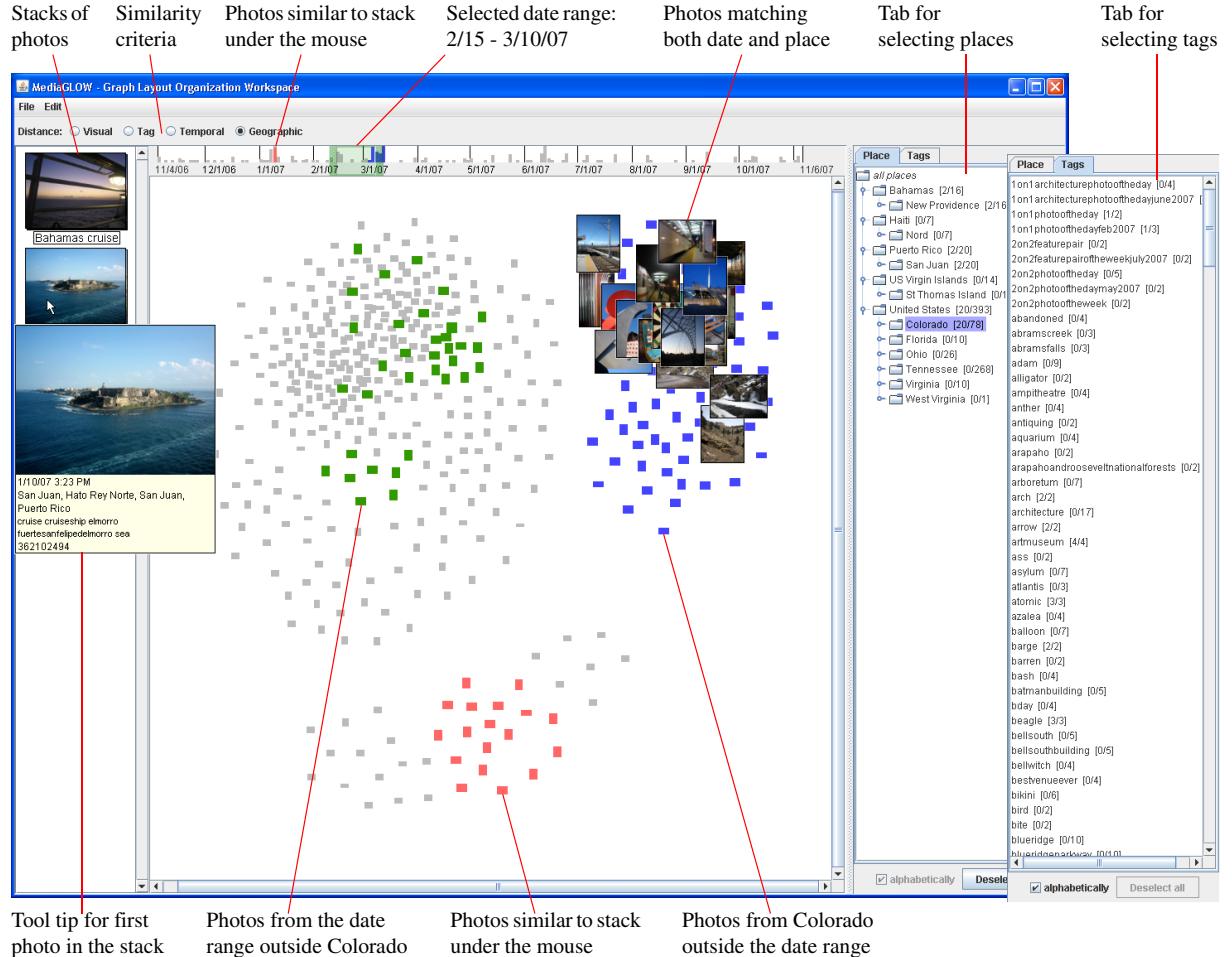


Figure 2.11: Photos Grouped by Geographic Similarity and Filtered by Date and Place.

MediaGLOW by Grgenohn et al. [16] is the application discussed in this paper. It's a content based image retrieval (CBIR) system with multiple ways to filter and sort the image collection.

The interface allows selecting a range of dates, places and tags at any time to filter the collection and the display will show the photos that match the filters, alongside indications of the existence of photos that match some of the filters. This display can then be arranged by four similarity criteria: temporal (by photo creation time), geographic (distances between places), tags (photos with similar tags are shown closer together) and visual (fig. 2.11).

The time is selected using a timeline on the top of the screen while tags and places are shown on the right side sorted alphabetically, by frequency of, in case of places, as a tree. Multiple selections are allowed to show more photos.

The photo display is graph based, allowing for overlapped images. Various metaphors were developed to ease the navigation of the collection. Zooming is allowed and changes both thumbnail positions and size for better experience, allowing the photos to spread away from each other but also increas-

ing the size so the user can have a better look at them. The authors think that bigger thumbnails and a correct grouping of related photos is more important than spreading them to avoid the overlapping problem.

Colour coding is used throughout the interface to help the user understand better what is being selected. For instance, on the timeline blue and grey are used to distinguish photos that match or not the selected location/tag. On the photo display, besides the photos that are actually shown are coloured blocks: blue for photos that match location only, green for time only and grey for those that didn't match anything.

Detailed user testing was performed and the importance of the multiple ways to organise and search the collections was emphasised since many systems are designed to have a single form of access. Some users also pointed the importance of being able to have a non overlapping view of the photos for part of the task.

Each view was found to have different levels of usage, the geographic being the most used and temporal the least, since it's very similar to the timeline. Visual similarity was less used than expected, even on collections where it was relevant.

2.13 Discussion

Currently there are a lot of approaches to image organisation and each has its own differences as demonstrated on Table 2.1.

Our survey revealed various methods for extracting the features of each image, from simple to complex.

One of the main problems is obtaining useful information from low level feature extraction of the image contents. Some try to get the most out of each image, with a variety of complex and time consuming procedures. Others try to focus on avoiding the complex computations by only getting simple but somewhat useful information. To contrast with this methods, Grgenohn's work [16] has found that users prefer other ways to filter the collection, like tags, dates and locations. It's still used, but probably isn't worth to spend much time on it with heavy processing.

Date and location are simple similarity measures and can be used to group the collections by events and locations like Naaman did on [25]. Current mainstream software like Apple's iPhoto¹ and Google's Picasa² are already doing it in a semi-automatic way.

Textual metadata like tags and descriptions are also widely used both on our survey and possibly on all major mainstream software. The problem with tags and descriptions is that people usually don't assign them to their photos but that's not a problem we're interested here.

The 3D Sphere from [34] is also interesting but doesn't provide a better interface to the collection than expanding the sphere surface grid view to a full screen grid view, keeping the zoom function. The Phorigami work [20] introduces some interesting metaphors for manipulating groups of photos, although

¹<http://www.apple.com/iphoto>

²<http://picasa.google.com>

Table 2.1: Different browsing methods

Work	Organisation				Visualisation	Focus	Size
	visual	date	gps	tags			
2.1 Qiu [28]	simple colour measures	—	—	—	Grid	Simplicity and Efficiency	60,000
2.2 Rodden [31]	—	—	—	✓	Grid	Similarity usefulness	100
2.3 Porta [27]	—	—	—	—	Spot (and others)	Unconventional visualisations	400
2.4 Cooper [12]	—	—	—	—	Static and animated	Usefulness of animations	—
2.5 Strong [36]	colour histograms and gradients	—	—	—	SOM	Evaluation of different features	2,200
2.6 Tan [38]	colour histograms of subdivisions	—	—	—	—	Evaluation of different features	12,000
2.7 Naaman [25]	—	✓	✓	—	—	Organization based on events	?
2.8 Chen [9]	colours, edges and textures	—	—	—	—	Efficient fast-sparse clustering	10,000
2.9 Heesch [18]	six different features	—	—	✓	Radial	Complex similarity network	32,000
2.10 Hsu [20]	—	—	—	—	Grid with groups	Interaction on grouped photos	1,333
2.11 Schaefer [34]	colour histograms of subdivisions	—	—	—	3D Sphere	3D mapping and interaction	4,500
2.12 Girgensohn [16]	—	✓	✓	✓	Overlapped graph	Having the best way to find photos	450

some clutter the view and could, therefore, be improved.

An interesting work is the Girgensohn's [16] visualisation approach, where images can be organised by various features and can be filtered down, displaying matched photos alongside placeholders for photos that are only match partially. It has some problems like image overlap and capacity for showing large collections.

Chapter 3

Solution Requirements

TODO: Fix below

Our survey was based on various types of previous work from the last ten years. Image browsers and technology have evolved a lot since then but there still isn't a definitive way for a user to look at its photo collection and understand its content and evolution.

We have the vision of a system that displays a large set of user's photos at the same time, in various arrangements, revealing patterns, differences and similarities between them.

For this to happen, the system must be able to handle tens of thousands of images and display them all on the screen while maintaining responsiveness. The system's UI should be clear and easy to use, allowing the user to navigate through the display of photos, through zooming and panning, and to reorganise the photos in a number of ways.

The system should also gather as much information as it can from the photos like, for instance, date and time, relevant colours, presence of people, type of photograph, location. While some of this information is already embedded in digital photographs as metadata, written by the digital camera when the photo was taken, others are usually not and need to be calculated or extracted. Faces and relevant colour information are an example of that and the system must be prepared to extract these features from the image. The system must allow other feature extraction methods to be easily added in the future. All this information will then be used by the user to reorganise and filter the photos on display.

We will now detail the work done, taking this requirements into account.

3.1 Motivation

Alguns casos de uso

3.2 Goals

"In this section we present the main goals to achieve in this work, as well as the design goals to the final solution described later in this document."

3.2.1 Main Goal

main goal

3.2.2 Design Goal

"Due to the different requirements of this work, we now introduce the several design goals that influenced its development."

Ease of Use

...

Chapter 4

Eagle Eye

Eagle Eye is a visualization tool that enables the display and manipulation of a large image set at once. It is focused at the regular computer user that has a few thousand digital photographs stored on the computer. It allows navigation, through panning and zooming the canvas of the image collection, sorting through different methods and multiple ways of filtering, either textual or visual.

We will now explain what makes all this work.

4.1 Architecture

Eagle Eye is divided in two parts: The visualization and the backend. As their names point, the former is what the user will manipulate to view the images while the latter reads the images and generates the needed data for the visualization to work.

We will detail this to parts next but first we will go through some design decisions that have shaped the work.

4.1.1 Design Decisions

DeepZoom

After analyzing some visualization technologies that allowed easy display and manipulation of images TODO: Insert a bunch of useless tech, Microsoft's Silverlight, with its DeepZoom technology, proved to be the best choice.

DeepZoom enables the use of multiple resolution images for efficient display at various zoom levels and, for instance, is used for the display of gigapixel photos TODO: links, where the user can view the whole image or zoom in the small details, or for collections of photos where the user can zoom between a view of multiple images and the details of a single image. So far, usages of DeepZoom have been restricted to promotion websites, art galleries and other closed usages.

This should be about the multi-scale images and not about DeepZoom → Our work aims at bringing this technology to the regular user, in a much simpler and dynamic way.

Although DeepZoom seemed a great technology, it relies on Silverlight, which by itself isn't as good as using a full-fledged desktop application framework like Windows Presentation Foundation (WPF) or other frameworks native to their platforms. This created some undesired limitations like the need to have two separated parts of the system, the visualization and the backend, and other smaller problems like limited access to the disk from the visualization part.

To create a DeepZoom application, some processing is required before hand to create multiple versions of each image in various resolutions and also to create imagery of a global view of the collection, also in multiple resolutions. This enables DeepZoom to only load the appropriate set of images for a certain zoom state,

4.1.2 Visualization

4.1.3 Backend

4.2 Backend

The backend is one of the two parts that make our system. It's responsible for reading the image files, processing all their information and generating the files that will be used by the visualization part.

We will now detail how all this works.

4.2.1 Structure

The backend is a command line utility that

4.3 Interface

Chapter 5

Evaluation

Chapter 6

Conclusions

Insert your chapter material here...

6.1 Achievements

The major achievements of the present work...

6.2 Future Work

A few ideas for future work...

Bibliography

- [1] ALBUZ, E., KOCALAR, E., AND KHOKHAR, A. A. Scalable color image indexing and retrieval using vector wavelets.
- [2] AZZAM, I., LEUNG, C., AND HORWOOD, J. A fuzzy expert system for concept-based image indexing and retrieval.
- [3] BEDERSON, B. B. Photomesa: a zoomable image browser using quantum treemaps and bubblemaps. pp. 71–80.
- [4] BERMAN, A., AND SHAPIRO, L. A flexible image database system for content-based retrieval. *Computer Vision and Image Understanding* (Jan 1999).
- [5] BRUIJN, O., AND SPENCE, R. Rapid serial visual presentation: a space-time trade-off in information presentation. *AVI '00: Proceedings of the working conference on Advanced visual interfaces* (May 2000).
- [6] CAMPBELL, I., AND VAN RIJSBERGEN, C. The ostensive model of developing information-needs. *CiteSeer* (Jan 2000).
- [7] CHANG, S., YAN, C., AND DIMITROFF, D. An intelligent image database system. *IEEE Transactions on ...* (Jan 1988).
- [8] CHEN, J.-Y., AND BOUMAN, C. A. Hierarchical browsing and search of large image databases. *Image Processing* (Jan 2002).
- [9] CHEN, J.-Y., BOUMAN, C. A., AND DALTON, J. C. Similarity pyramids for browsing and organization of large image databases. *Human Vision and Electronic Imaging III* 3299 (Jan 1998), 563–575.
- [10] CHRISTMANN, O., AND CARBONELL, N. Browsing through 3d representations of unstructured picture collections: an empirical study. *Proceedings of the working conference on Advanced visual interfaces* (2006), 445–448.
- [11] COMBS, T., AND BEDERSON, B. Does zooming improve image browsing? ... of the fourth ACM conference on ... (Jan 1999).
- [12] COOPER, K., BRUIJN, O., SPENCE, R., AND WITKOWSKI, M. A comparison of static and moving presentation modes for image collections. *AVI '06: Proceedings of the working conference on Advanced visual interfaces* (May 2006).

- [13] CUNNINGHAM, S., AND MASOODIAN, M. Identifying personal photo digital library features. 401.
- [14] DATTA, R., JOSHI, D., LI, J., AND WANG, J. Image retrieval: Ideas, influences, and trends of the new age. *ACM Computing Surveys (CSUR)* (Jan 2008).
- [15] DENG, J., DONG, W., SOCHER, R., LI, K., AND FEI-FEI, L. Imagenet: A large-scale hierarchical image database. *computer.org* (Jan 2009).
- [16] GIRGENSOHN, A., SHIPMAN, F., TURNER, T., AND WILCOX, L. Flexible access to photo libraries via time place, tags, and visual features. pp. 187–196.
- [17] HEESCH, D. A survey of browsing models for content based image retrieval. *Multimedia Tools and Applications* (Jan 2008).
- [18] HEESCH, D., AND RÜGER, S. Nnk networks for content-based image retrieval. *Advances in Information Retrieval* (Jan 2004), 253–266.
- [19] HILLIGES, O., BAUR, D., AND BUTZ, A. Photohelix: Browsing, sorting and sharing digital photo collections. *Horizontal Interactive Human-Computer Systems, International Workshop on 0* (2007), 87–94.
- [20] HSU, S., CUBAUD, P., AND JUMPERTZ, S. Phorigami: A photo browser based on meta-categorization and origami visualization. *Human-Computer Interaction. Novel Interaction Methods and Techniques* (2009), 801–810.
- [21] JAIN, A., AND VAILAYA, A. Image retrieval using color and shape. *Pattern Recognition* (Jan 1996).
- [22] KRISHNAMACHARI, S., AND ABDEL-MOTTALEB, M. Image browsing using hierarchical clustering. *Computers and Communications, 1999. Proceedings. IEEE International Symposium on* (1999), 301 – 307.
- [23] MINKA, T., AND PICARD, R. An image database browser that learns from user interaction. *Master's thesis* (Jan 1996).
- [24] MINKA, T., AND PICARD, R. Interactive learning with a “society of models”. *Computer Vision and Pattern Recognition, 1996. Proceedings CVPR '96, 1996 IEEE Computer Society Conference on* (1996), 447 – 452.
- [25] NAAMAN, M., SONG, Y., AND PAEPCKE, A. Automatic organization for digital photographs with geographic coordinates. *Digital Libraries* (Jan 2004).
- [26] PLATT, J., AND CZEKALA, M. Phototoc: Automatic clustering for browsing personal photographs. ... *Communications and Signal ...* (Jan 2004).
- [27] PORTA, M. Browsing large collections of images through unconventional visualization techniques. *AVI '06: Proceedings of the working conference on Advanced visual interfaces* (May 2006).

- [28] QIU, G., MORRIS, J., AND FAN, X. Visual guided navigation for image retrieval. *Pattern Recognition* 40, 6 (2007), 1711–1721.
- [29] QIU, G., YE, L., AND FENG, X. Fast image indexing and visual guided browsing. . . . *Workshop on Content-Based Multimedia Indexing* (2010).
- [30] RODDEN, K. How do people organise their photographs. *BCS IRSG 21st Ann. Colloq. on Info. Retrieval* . . . (Jan 1999).
- [31] RODDEN, K., BASALAJ, W., SINCLAIR, D., AND WOOD, K. R. Does organisation by similarity assist image browsing? *Proceedings of the SIGCHI conference on Human factors in computing systems* (2001), 197.
- [32] RODDEN, K., AND WOOD, K. R. How do people manage their digital photographs? *Proceedings of the SIGCHI conference on* . . . (Jan 2003).
- [33] RUI, Y., HUANG, T., AND CHANG, S. Image retrieval: Current techniques, promising directions, and open issues* 1. *Journal of visual communication and image* . . . (Jan 1999).
- [34] SCHAEFER, G. A next generation browsing environment for large image repositories. *Multimedia Tools and Applications* (Jan 2010).
- [35] STRONG, G., AND GONG, M. Browsing a large collection of community photos based on similarity on gpu. *ISVC '08: Proceedings of the 4th International Symposium on Advances in Visual Computing, Part II* (Dec 2008).
- [36] STRONG, G., AND GONG, M. Organizing and browsing photos using different feature vectors and their evaluations. *CIVR '09: Proceeding of the ACM International Conference on Image and Video Retrieval* (Jul 2009).
- [37] TAMURA, H., AND YOKOYA, N. Image database systems: A survey. *Pattern Recognition* (Oct 2002). Extensa lista de image browsers....
- [38] TAN, K., OOI, B., AND YEE, C. An evaluation of color-spatial retrieval techniques for large image databases. *Multimedia Tools and Applications* (Jan 2001).
- [39] TORRALBA, A., FERGUS, R., AND FREEMAN, W. T. 80 million tiny images: A large data set for nonparametric object and scene recognition. *IEEE Trans. Pattern Anal. Mach. Intell.* 30, 11 (2008), 1958–1970.
- [40] WORRING, M. Interactive access to large image collections using similarity-based visualization. *Journal of Visual Languages & Computing* (Jan 2008).

