



## EagleEye

Large photoset visualization

**Carlos Eduardo Henriques de Jesus Fonseca**

Dissertação para a obtenção de Grau de Mestre em  
**Engenharia Informática e de Computadores**

### Júri

Presidente: João Emilio Segurado Pavão Martins  
Orientador: Daniel Jorge Viegas Goncalves  
Vogal: Carlos Martinho

Outubro 2011



## **Acknowledgments**

I want to thank my thesis adviser, Prof. Daniel Gonçalves for his support, vision and guidance that helped me to accomplish this work.

I also thank my family for all of their support, patience and for bearing with me through this long path.

A very special thanks to my dearest friends who have been with me through all this time, both providing good and serious moments that have shaped a better person out of me.



# Resumo e palavras chave

A fotografia digital tem feito parte da vida das pessoas na ultima década, crescendo nos discos rígidos, mas a visualização de imagens não tem vindo a evoluir o suficiente para acompanhar este crescimento. A maior parte das pessoas guarda as fotos em pastas ou utiliza programas que não conseguem mostrar grandes quantidades de fotografias ao mesmo tempo, não fornecendo uma visão global da biblioteca.

Este trabalho cresceu das características de outros trabalhos de visualizações alternativas de imagens e tenta trazer a utilizadores comuns uma melhor maneira de explorar as suas bibliotecas de fotografias e aprender mais sobre elas.

Este trabalho fornece um sistema extensível para a extracção e agregação de características das imagens do utilizador e respectivos meta-dados, Actualmente são extraídas cores, rostos, datas, pastas e palavras-chave, mas, já que é extensível, poderá ser utilizado para muito mais.

Também criámos uma aplicação de visualização que usa os dados processados e os disponibiliza aos utilizadores numa interface limpa e simples. A aplicação começa por apresentar no ecrã, numa grelha, todas as imagens adicionadas ao programa, agrupados por datas. A partir daqui, o utilizador pode alterar a visualização ampliando-a ou reduzindo-a, organizando as imagens por cores, rostos ou pastas, podendo ainda filtrar utilizando uma mistura de todos os dados recolhidos pelo sistema de processamento. Desta forma, os utilizadores podem visualizar e explorar suas bibliotecas de uma forma diferente daquilo a que estão habituados.

**Palavras-chave:** fotografias, visualização de imagens, exploração de imagens, extração de características



# Abstract and keywords

Digital photography has been a part of people's lives for the past decade, growing on hard drives, but image visualization has not been evolving enough to accompany this growth. Most people keep photos on folders or use software that can't display a large number of photos at the same time, not providing a good overview of their collection.

This work grew from the features of other works of alternative image visualizations and tries to bring a better way for regular people to explore their photo libraries and learn more about them.

This work provides an extensible backend system intended for image and metadata processing, gathering information about the user's photos, currently colors, faces, dates, paths and keywords, but, since it's extensible, could be used for much more.

We also created a visualization application to use the processed data and provide it to the users, in an interface that is clean and simple. It starts by displaying all the added images on the screen, at the same time, in a grid, grouped by dates. From here, the user can zoom in and out and change the display to organize the images by colors, faces or paths, and can filter using a mix of any data gathered by the backend . This way, users can view and explore their libraries in a different way from what they have been doing.

**Keywords:** photographs, image visualization, image browsing, feature extraction



# Contents

Acknowledgments . . . . .	iii
Resumo e palavras chave . . . . .	v
Abstract and keywords . . . . .	vii
List of Tables . . . . .	xi
List of Figures . . . . .	xiv
Acronyms . . . . .	xv
<b>1 Introduction</b>	<b>1</b>
1.1 Our vision . . . . .	2
1.2 Contributions . . . . .	2
1.3 Structure of the Document . . . . .	3
<b>2 Related Work</b>	<b>5</b>
2.1 Related work . . . . .	5
2.1.1 Visual guided navigation for image retrieval . . . . .	5
2.1.2 Does organization by similarity assist image browsing? . . . . .	6
2.1.3 Browsing large collections of images through unconventional visualization techniques . . . . .	6
2.1.4 A comparison of static and moving presentation modes for image collections . . . . .	7
2.1.5 Organizing and browsing photos using different feature vectors and their evaluations . . . . .	7
2.1.6 An evaluation of color-spatial retrieval techniques for large image databases . . . . .	9
2.1.7 Automatic organization for digital photographs with geographic coordinates . . . . .	9
2.1.8 Similarity pyramids for browsing and organization of large image databases . . . . .	9
2.1.9 NN <sup>k</sup> networks for content-based image retrieval . . . . .	10
2.1.10 Phorigami: A Photo browser based on meta-categorization and origami visualization	11
2.1.11 A next generation browsing environment for large image repositories . . . . .	12
2.1.12 Flexible access to photo libraries via time place, tags, and visual features . . . . .	13
2.1.13 PhotoMesa: a zoomable image browser using quantum treemaps and bubblemaps	14
2.2 Discussion . . . . .	15
<b>3 Solution Requirements</b>	<b>19</b>
3.1 Main Goals . . . . .	19

3.2 Implementation Requirements . . . . .	20
3.3 Features of Images and Photographs . . . . .	21
<b>4 Eagle Eye</b>	<b>25</b>
4.1 Overview . . . . .	25
4.2 Backend . . . . .	25
4.2.1 Architecture of the Backend . . . . .	26
4.2.2 Feature Extraction Plugins . . . . .	28
4.3 Visualization . . . . .	32
4.3.1 Overview . . . . .	32
4.3.2 Disposition of Images on Canvas . . . . .	35
4.3.3 Architecture of the Visualization . . . . .	38
<b>5 Evaluation</b>	<b>43</b>
5.1 Objectives . . . . .	43
5.2 Testing Methodology . . . . .	43
5.3 Characterization of users . . . . .	44
5.4 Test Results . . . . .	44
5.4.1 User behavior during exploration . . . . .	45
5.4.2 Results for tasks . . . . .	45
5.4.3 Errors and misunderstandings . . . . .	46
5.4.4 System Usability Scale . . . . .	48
5.4.5 What users liked . . . . .	48
5.4.6 User suggested improvements . . . . .	49
5.5 Accomplished Objectives . . . . .	50
5.6 Case Studies . . . . .	51
5.7 Discussion . . . . .	52
<b>6 Conclusions</b>	<b>53</b>
6.1 Conclusion . . . . .	53
6.2 Future Work . . . . .	54
6.2.1 Consolidation . . . . .	54
6.2.2 Feature Extraction Plugins . . . . .	54
6.2.3 Visualization . . . . .	54
6.2.4 Connect . . . . .	55
<b>A User Survey</b>	<b>57</b>
A.1 User characterization . . . . .	57
A.2 Characterization of Photo Library . . . . .	58
<b>Bibliography</b>	<b>65</b>

# List of Tables

2.1 Different browsing methods . . . . .	16
3.1 Example of some relevant EXIF fields imprinted by the camera and by the computer software used for post processing. . . . .	23
5.1 Average of the results for each of the ten questions of System Usability Scale (SUS). . . . .	48



# List of Figures

2.1 Chromacy diagram and color related images, from the work of Qiu et al. . . . .	5
2.2 Three arrangements of 100 images of Kenya, based on visual similarity, from the work of Rodden and Sinclair . . . . .	6
2.3 The two better visualizations from Porta's work. . . . .	7
2.4 The six rapid serial visual presentation modes used in the experiments . . . . .	8
2.5 The result obtained for organizing 2200+ photos using color autocorrelogram feature vector, using Strong's work [55]. . . . .	8
2.6 Images organized with [12]. Images with similar color and texture are spatially adjacent. .	10
2.7 Local network around the chosen butterfly image depicted in the middle. . . . .	11
2.8 Hsu's work of grouping together related images . . . . .	12
2.9 Hue sphere of a dataset, from Schaefer's work. . . . .	12
2.10 Photos grouped by geographic similarity and filtered by date and place. . . . .	13
2.11 PhotoMesa with over 500 images in 17 groups. . . . .	15
4.1 Basic architecture of Eagle Eye's backend . . . . .	26
4.2 Process in use by our system for importing a folder using ExifTool. . . . .	27
4.3 Example of information extracted from an image. . . . .	27
4.4 A colorful photo. . . . .	28
4.5 Comparison of full color and reduced color images. . . . .	29
4.6 Color palette in use for restricting the possible colors in an image. . . . .	29
4.7 Comparison between the original image, the adaptive color reduction method and the color mapping method. . . . .	30
4.8 Example of the OpenCV library detecting faces on a common photo. . . . .	31
4.9 Detection vs. speed results of the face detection test. . . . .	31
4.10 Example of the visualization of 5679 images, organized by capture date, in Eagle Eye. .	33
4.11 The toolbar on top of the visualization user interface (UI). . . . .	34
4.12 Overlays in Eagle Eye . . . . .	34
4.13 Division of the hue values in 12 equal parts, 30 points apart from each other. . . . .	35
4.14 Common treemap based on space filling. . . . .	36
4.15 Example of the linear display. . . . .	37
4.16 Examples of filter suggestions. . . . .	38

4.17 Selected examples of photo stacking as presented by Eagle Eye. . . . .	40
4.18 Real world example of stacking on a trip's photos with HDR and Panorama images. . . . .	41
5.1 User characterization. . . . .	45
5.2 Graph of time spent on each task . . . . .	47
5.3 Graph of errors made on each task . . . . .	47
5.4 Case study: six time lapses of 563 images, reduced to 6 spaces on the canvas . . . . .	51
A.1 Kinds of photographers. . . . .	57
A.2 Library Size . . . . .	58
A.3 Software used . . . . .	58
A.4 Systematic organization in folders, in an application or no organization. . . . .	59
A.5 Users that take photos in close sequence. . . . .	59
A.6 User rated features. . . . .	59
A.7 Whether people think they can get a good overview or not. . . . .	60

# Acronyms

**UI** user interface

**UX** user experience

**SOM** self organizing map

**CBIR** content based image retrieval

**MP** megapixel

**FEP** feature extraction plugin

**WPF** Windows Presentation Foundation

**EXIF** Exchangeable image file format

**HDR** High Dynamic Range

**SUS** System Usability Scale



# Chapter 1

## Introduction

***photography***

noun

the art or practice of taking and processing photographs.

---

New Oxford American Dictionary 3rd edition

Photography is a very personal activity. Each person has its own way of doing it, with its own devices, software and techniques. It ranges from common people who take photos of their cats using their low quality camera phones to put on the internet, up to professional photographers with big cameras and lenses to take award-winning photos and everything in between. Digital cameras have brought photography to the masses with its ease of use, no cost per photo and easy processing made photography skyrocket. Flickr<sup>1</sup> has reached the 6.000.000.000 upload in August 2011 and seeing a 20% increase year-over-year since the website's debut in 2004<sup>2</sup>.

Not only are people taking more and more photos, some photography techniques that use multiple photo captures also appeared or, at least, became more relevant with the digital cameras among some enthusiasts and professional photographers. Examples of this are burst photographs, where many photos are taken in a quick sequence to capture a motion sequence, or panoramas that are made from various overlapping photos of a subject that is bigger than what the camera can capture or even something called High Dynamic Range (HDR) photography that merges multiple captures to correct the lack of capacity of the sensors to capture very dark and very bright areas on the same scene. Yet another technique, this one more related to video, is the production of time lapses, large sequences of photographs taken with a fixed interval between them, capturing life at a slower rate. Time lapse photos are usually put together in a video, making hours pass by in seconds. We did a survey and found out 40% does some of this kinds of photography.

This growth of the digital brings the problem that photo collections have started to grow more than they used to, in the film era. But storing digital photographs isn't that different from the past. In fact, people have been storing their photographs on folders on their computers instead of photo albums on

---

<sup>1</sup><http://flickr.com> is, probably, the most used photo-sharing website worldwide.

<sup>2</sup>According to a blog post on the Flickr website <http://blog.flickr.net/en/2011/08/04/6000000000>

the shelf. One could say computers can help people view, explore and find more photos in a shorter period of time, they certainly have the power for that but, in the end, they usually don't do much more than make the user look for the photo album (probably a folder or an "album" on some applications) and then flick through its pages (scrolling) until the photos the user was looking for appear.

This is true for the most common photo management software on the market, like Google's Picasa<sup>3</sup>, Apple's iPhoto<sup>4</sup> or Aperture<sup>5</sup>, or Adobe's Lightroom<sup>6</sup>. Although they bring some management improvements with them, the analogy above still applies. They all provide a way to select albums or folders and scroll through contents. They allow searching through some existing metadata, the only metadata that some of them generate is only face detection and end up having lots of buttons, toggles and options for their editing capabilities that clutter the interface. All of this, added to the fact that collections quickly reach the thousands of images, prohibit a global view of the collection. It's not easy to gather all the images that are spread across the system and view them all together, understanding its evolution and characteristics.

## 1.1 Our vision

We want to provide users a totally different way to visualize their photos by focusing only on them.

We propose that visualizing photographs in a birds-eye view, with a number of different ways to organize, group, sort and filter them by using features like events or dates, dominant color or people in photos will provide users a much more satisfying experience of discovery of their libraries.

To demonstrate this, we implemented Eagle Eye, a system that allows users to do just that. After testing, we observed that our approach allows users to navigate and explore their photo libraries in a whole new way, revealing patterns or evolution over time, while enjoying their photos.

## 1.2 Contributions

We created a way to bring image exploration to the regular users, with their own photos.

We built a flexible and extensible backend system to extract interesting information from users' images. We have implemented a few plugins to extract some features, from metadata to median color and face detection, and to generate visualization data. It's possible to change the whole visualization and only change one plugin to adapt to the new system.

We also took an existing technology that allows image zooming and made it easy for regular people to use it with their own images and have a nice experience using it, with various complex features simplified. Users can use it to view, sort, group and filter images in their photo library.

---

<sup>3</sup><http://picasa.google.com>

<sup>4</sup><http://www.apple.com/iphoto>

<sup>5</sup><http://www.apple.com/aperture>

<sup>6</sup><http://www.adobe.com/lightroom>

### **1.3 Structure of the Document**

We will start by going through some previous work related to ours (chapter 2) followed by an identification of the requirements for our work (chapter 3). Next, we will see what was implemented and how it works (chapter 4) and how the users responded to it (chapter 5). To conclude, we will discuss some work that could be done to improve this thesis and some final thoughts on the work done (chapter 6).



# Chapter 2

## Related Work

Interactive image visualization techniques have been explored for some time now, many of them related to image organization or retrieval in large collections. There have been some interesting ideas across the board and we will now take a look at some of them.

### 2.1 Related work

#### 2.1.1 Visual guided navigation for image retrieval

Qiu et al. [48] explore the requirements of a system intended for visualizing large photo collections. They identify as the two most important requirements, the first being an easy to use UI, that gives clean information to the user and helps to create a mental image of the whole collection helping him to navigate on the collection. The second requirement is responsiveness because while image processing can be an heavy task, the user needs to be able to interact with the interface and he won't use the application if it's slow.

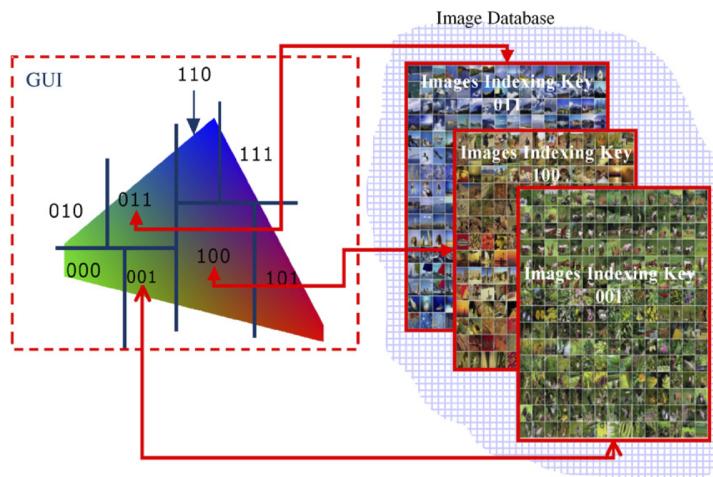


Figure 2.1: The chromacy diagram is split in parts and each image belongs to one of this parts. The diagram is part of the UI and when navigating through the diagram, only the images related to that part are shown.

The system shows all the photos arranged by color, just as many others like it. The difference is the process in use. Instead of calculating distance vectors based on the histogram of each image, this approach classifies each image with a simple description, like an average of its colors, and arranges them by that value, on a color map (fig. 2.1). The process is much faster, but is also more error prone, specially on photos without a clear main color.

Their tests show they achieved good responsiveness and better results than using a file explorer.

### 2.1.2 Does organization by similarity assist image browsing?



Figure 2.2: Three arrangements of 100 images of Kenya, based on visual similarity. On the left is the arrangement with overlap, in the middle a 12x12 grid (which removes the overlap while preserving some of the structure), and on the right a 10x10 grid (which maximises the thumbnail size).

The aim of this work by Rodden and Sinclair [51] was to evaluate how photo organization by similarity (fig. 2.2) could benefit a user looking for images. Some users tested an application that could show the same images both in a random and in an organized by similarity way. This organization by similarity was based on a rough description of the images, but it could be other descriptors.

The results differ if the user knows what he's looking for or not. In case he does, being able to filter only the relevant images makes it quick to find the ones that matter. This obviously depends on the quality of the labeling. Users reported that sometimes the similar images appear to merge.

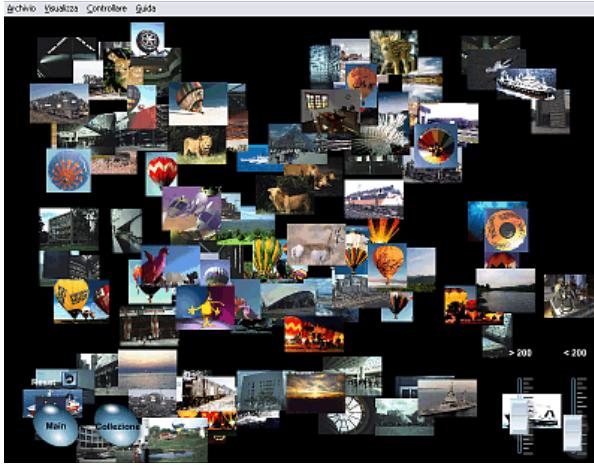
In case the user doesn't know what he's looking for, the random approach might be helpful because the strong images usually contrast to their neighbors and thus appear to stand out.

For some people, having access to different arrangements of the same set of images is useful, although the source of the individual differences still needs to be determined.

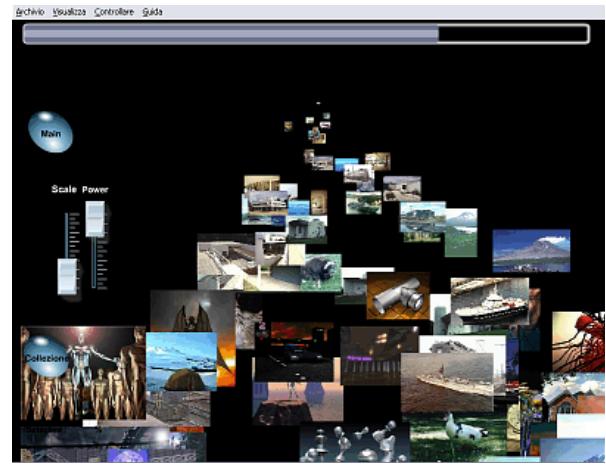
### 2.1.3 Browsing large collections of images through unconventional visualization techniques

Porta [47] describes a few visualization methods for large collections of images and tests them with users. The purpose is to find ways or metaphors that provide a good visualization experience in terms of time spent and quality of the visualization.

Some of the various techniques were the simple image grid view, a grid view with variable and



2.3.1: Spot display



2.3.2: Shot display

Figure 2.3: The two better visualizations from Porta's work.

independent height and width (EIB), a view that animates images like they were shot from a distance and get closer to the user called Shot (fig. 2.3.2), a view where images quickly appear on random positions on screen named Spot (fig. 2.3.1), and some other less commons like one that simulates a cylinder created with images (Cylinder), and others less relevant (Rotor, Tornado and Tornado of Planes).

The testing was based on the efficiency of users searching for specific images on a collection of a thousand photos. The Spot view was the best, followed the Shot, Cylinder finally the common grid view. The other views got scores near or below the grid view.

#### 2.1.4 A comparison of static and moving presentation modes for image collections

This paper by Cooper et al. [16] is not about large libraries, but about what kind of interfaces for image showing has greater success of user recognition and preference (fig. 2.4).

With the help of eye-tracking techniques and user preference, the authors determined that static images are better than moving ones because makes them easier to recognize and avoid some user confusion.

#### 2.1.5 Organizing and browsing photos using different feature vectors and their evaluations

Although it doesn't mention why, Strong [55] focuses on the better experience provided by color organization of a large image collection (fig. 2.5).

A self organizing map (SOM) is used to display the images on the screen featuring zooming, panning and sorting capabilities. The work is then based around the various methods used to determine the images' similarity.

Simpler methods are based on color histograms, which aren't affected by rotations or scales but, by not having spacial information on colors, allows very different images be closer together.

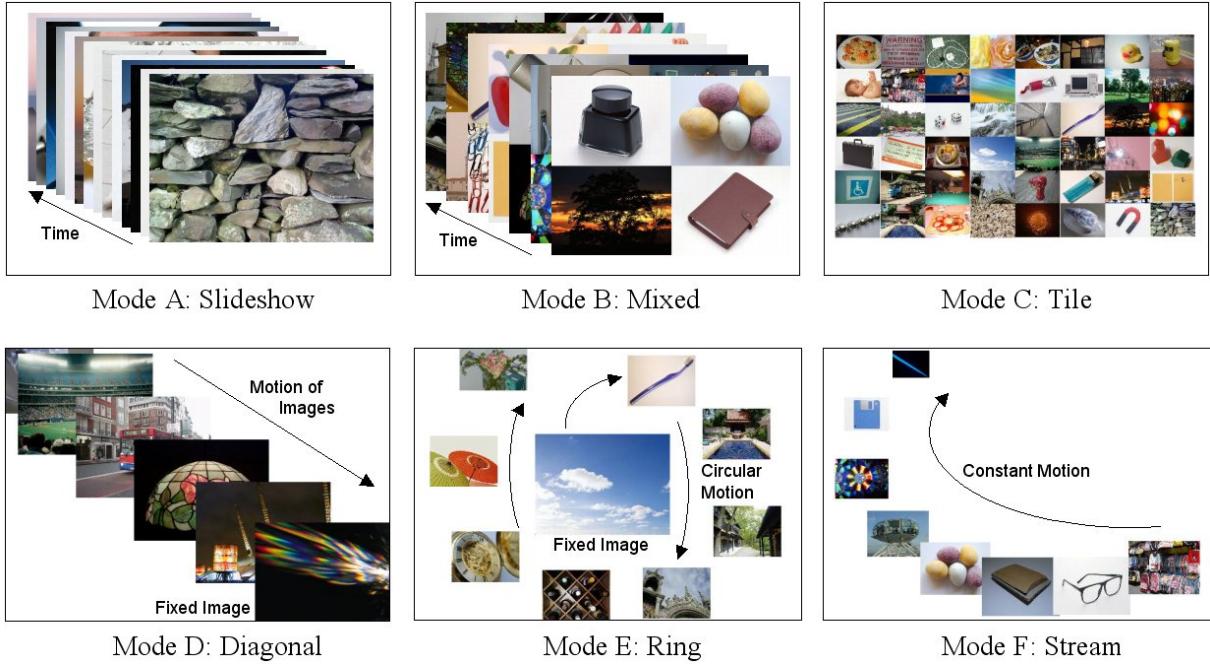


Figure 2.4: The six rapid serial visual presentation modes used in the experiments

Other methods rely on gradients which contain spatial information and, therefore, are sensitive to image contents, but not color. In general, the best methods were found to be the hybrid ones, where both color histograms and gradients are used to classify the images. No user testing is made in this project, neither is the speed of image categorization referred about the methods used.



Figure 2.5: The result obtained for organizing 2200+ photos using color autocorrelogram feature vector, using Strong's work [55].

## **2.1.6 An evaluation of color-spatial retrieval techniques for large image databases**

Tan et al. [58] present an evaluation of three color-spatial image retrieval techniques.

The signature-based technique creates a signature for each image, based on the most frequent colors, according to a threshold, of each subdivision, or bin, of that image. The comparison between images is made by comparing the main colors present on each bin. It is possible to assign more weight to specific bins according to the user's interest.

The partition-based approach is also based on bins, each having its own color histogram. The similarity between images is given by the distance of the histograms of the corresponding bins.

The cluster-based method bases on the fact that humans focus on large patches (clusters) of the same color and, therefore, two images will appear similar if both have similar colored clusters on at roughly the same location. This method extracts the larger clusters and their color from each image. The similarity is calculated by the amount of overlap between clusters.

This techniques were tested with a collection of 12,000 images and, besides the color information, the brightness was also analyzed for increased performance. The authors conclude the signature method was generally better on both effectiveness and efficiency.

## **2.1.7 Automatic organization for digital photographs with geographic coordinates**

In this paper, by Naaman et al. [43], is described a system that organizes digital photographs accordingly to location and date embedded on the metadata.

The objective was trying to mimic the way people think about their collections. Photos are usually bursts separated by some time. Based on this and on the different places, events can be created to agglomerate photos from the same bursts. Location naming is done by calculating the most relevant places, like parks or cities, and then mixing the more precise locations with the more important neighbor cities to create a relevant and identifiable name. This was specially important since this work didn't involve showing any maps, but only the location names and events.

The authors showed good results and, nowadays, some common applications use similar features although including maps.

## **2.1.8 Similarity pyramids for browsing and organization of large image databases**

Chen et al. [12] present a method for designing a hierarchical browsing environment called a similarity pyramid. The similarity pyramid groups similar images together while allowing users to view the database at varying levels of resolution. Each level is organized such that similar images are in close proximity on a two-dimensional grid (fig. 2.6). Images are first organized into a binary tree through agglomerative clustering based on color, edge and texture similarities. The binary tree is transformed into a quadtree, a tree in which each node has four children instead of only two.

The similarity pyramid is best constructed using agglomerative (bottom-up) clustering methods, and

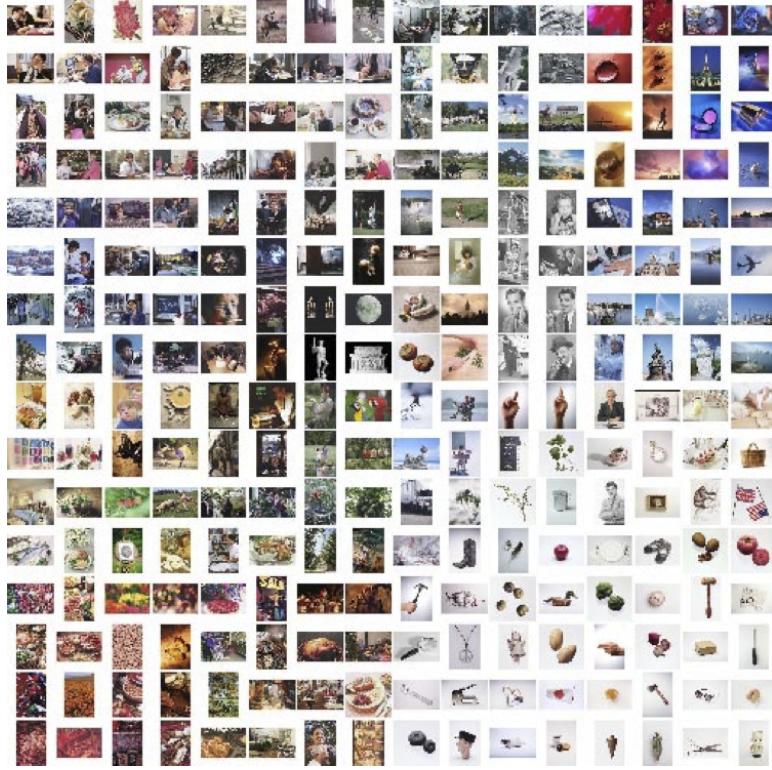


Figure 2.6: Images organized with [12]. Images with similar color and texture are spatially adjacent.

a fast-sparse clustering method is presented which dramatically reduces both memory and computation over conventional methods. This method is based on the flexible agglomerative clustering algorithm, but using only a sparse proximity matrix and exploiting the author's approximate branch and bound search algorithm.

The authors found that the method for mapping the clustering to a pyramid can make a substantial difference in the quality of organization. Finally, a dispersion metric for objectively measuring pyramid organization was proposed, and found that it correlated well with the author's subjective evaluations of pyramid organization.

### 2.1.9 $NN^k$ networks for content-based image retrieval

Heesch [30] describes a different interaction technique for content based search in large image collections. Each image is a vertex in a graph and arcs are established between images if there exists at least one combination of features for which one image is retrieved as the nearest neighbor of the other. Each arc is weighted by the proportion of feature combinations for which the nearest neighbor relationship holds. By integrating the retrieval results over several feature combinations, the resulting network helps expose the semantic richness of images.

The interface reflects the vertexes and respective arcs, allowing to browse between the related images (fig. 2.7) in the network.

Seven low-level features are used for the classification of the images. HSV Global color Histograms maps the images by color, saturation and brightness; color Structure Descriptor maps the distribution

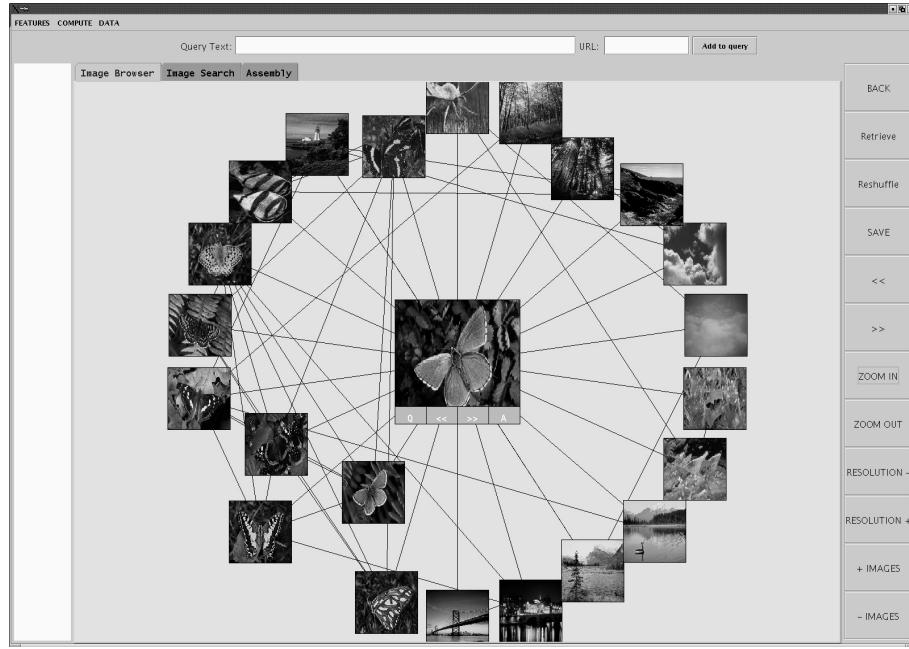


Figure 2.7: Local network around the chosen butterfly image depicted in the middle.

of colors by dividing each image in 64 windows, the color space in 184 bins and associating color bins with image windows; Thumbnail feature compares identical images by saving the grey value of each pixel from a scaled down version of the image; Convolution filters discovers very selective features by reapplying 25 filters three times; Variance feature calculates standard deviations within a sliding window; Uniformity Feature is another statistical feature, calculating the grey level of an image split in 64 parts; Bag of words is the last feature and weights words associated to each image.

Tests showed great results when using a mix of search, relevance feedback and browsing, and even browsing by itself was considerably better than other, more restrictive, approaches.

The technique helps avoid the problem of image polysemy by showing all gathered meanings of the images to the user. The feature network is pre-computed, allowing for quick realtime browsing. The authors claim it took 50 hours to process 32000 images, but made no reference to the possibility adding images to the collection, after the computation.

### 2.1.10 Phorigami: A Photo browser based on meta-categorization and origami visualization

Hsu et al. [34] try to ease the browsing problem by analyzing the collections and identifying groups of related pictures. Each type of group is visualized in a specific way, inspired by the Origami art.

Groups of similar or related photos were manually classified based on camera movement and subject movement, creating different types of groups static view where both camera and subject are fixed and is presented as a panorama; multi-view where the subject is fixed, but the camera is moving and is shown as a presentation; if the subject is moving, the photos are categorized as motion capture and can be shown as an animated photo (fixed camera) or a presentation (moving camera); finally group photos,

where different groups of people are photographed, are shown as a folding presentation.

This covers various cases where the photographer takes a few similar photos of the same subject because it's either a panorama, various angles or just to be sure the photo was well captured.

The interface implements the different presentation types as different metaphors, easy for the user to understand, like a folded paper on a wide panorama that can be expanded (fig. 2.8). Although some of them appear to be a little hard to distinguish in its compressed form, it shouldn't be difficult to make it clearer. Other possible problem is the use of different touch interactions for each presentation type that might confuse users on what gesture should they use.



Figure 2.8: On the left, an example of an interaction on a group of photos that makes a panorama. On the right, a visualization on 537 photos with some groups.

### 2.1.11 A next generation browsing environment for large image repositories



Figure 2.9: Hue sphere of a dataset, from Schaefer's work.

Schaefer [54] tries to take similarity based organization of images from the 2D space to a 3D sphere, which allows interaction from the users. Rotating the sphere reveals images with different colors while tilting it reveals brighter or darker images.

Large image collections are handled through a hierarchical approach that brings up similar, previously hidden, images when zooming in on an area.

The description of the color is retrieved by calculating the image's median color for its efficiency over other methods like histograms. This two features are directly mapped onto the sphere's coordinates and

the entire structure is pre-calculated so browsing can be performed in real-time. Image overlapping is also avoided (fig. 2.9).

The work was tested on a 4500 image collection with no evaluation as to its performance and a weak and subjective user testing.

### 2.1.12 Flexible access to photo libraries via time place, tags, and visual features

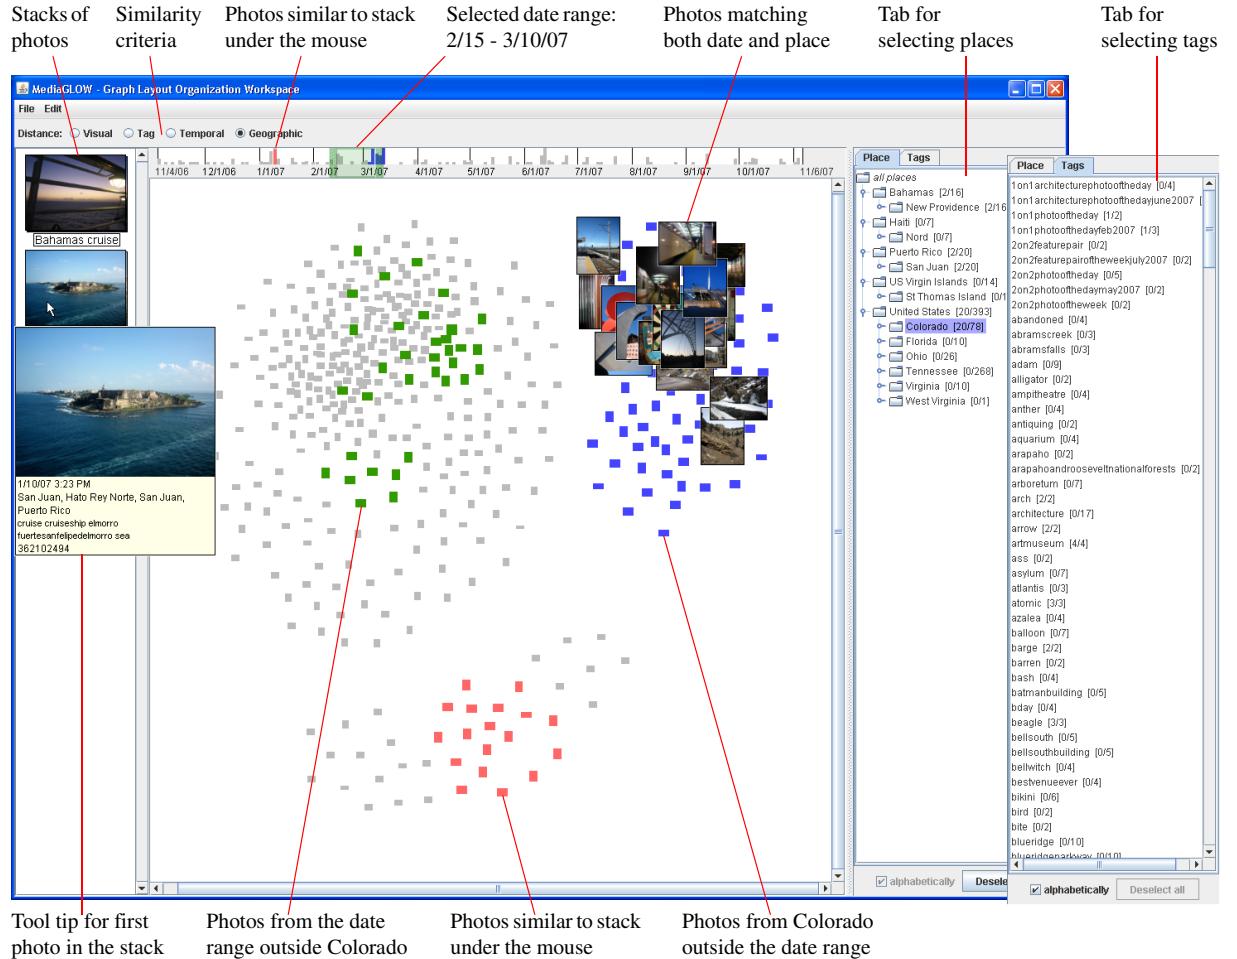


Figure 2.10: Photos grouped by geographic similarity and filtered by date and place.

MediaGLOW by Grgenohn et al. [27] is the application discussed in this paper. It's a content based image retrieval (CBIR) system with multiple ways to filter and sort the image collection.

The interface allows selecting a range of dates, places and tags at any time to filter the collection and the display will show the photos that match the filters, alongside indications of the existence of photos that match some of the filters. This display can then be arranged by four similarity criteria: temporal (by photo creation time), geographic (distances between places), tags (photos with similar tags are shown closer together) and visual (fig. 2.10).

The time is selected using a timeline on the top of the screen while tags and places are shown on the right side sorted alphabetically, by frequency of, in case of places, as a tree. Multiple selections are allowed to show more photos.

The photo display is graph based, allowing for overlapped images. Various metaphors were developed to ease the navigation of the collection. Zooming is allowed and changes both thumbnail positions and size for better experience, allowing the photos to spread away from each other, but also increasing the size so the user can have a better look at them. The authors think that bigger thumbnails and a correct grouping of related photos is more important than spreading them to avoid the overlapping problem.

Color coding is used throughout the interface to help the user understand better what is being selected. For instance, on the timeline blue and grey are used to distinguish photos that match or not the selected location/tag. On the photo display, besides the photos that are actually shown are colored blocks: blue for photos that match location only, green for time only and grey for those that didn't match anything.

Detailed user testing was performed and the importance of the multiple ways to organize and search the collections was emphasized since many systems are designed to have a single form of access. Some users also pointed the importance of being able to have a non overlapping view of the photos for part of the task.

Each view was found to have different levels of usage, the geographic being the most used and temporal the least, since it's very similar to the timeline. Visual similarity was less used than expected, even on collections where it was relevant.

### **2.1.13 PhotoMesa: a zoomable image browser using quantum treemaps and bubblemaps**

This work presents PhotoMesa by B. B. Bederson [3], an application that supports browsing sets of images in a zoomable environment. It also supports clustering by metadata, not requiring previous work from the user. Users can choose directories of images and they are all displayed in a space-filling manner fig. 2.11. Images are displayed in groups, from the directories they origin from and users can smoothly zoom in to a group and then to a single image. It generates multiple-resolution thumbnails for maintaining a good performance. Keyboard keys are also used to navigate the canvas. The groups display their name and have different background colors for a better distinction. It also supports drag and drop of images to other applications. Text search is possible as well as selecting a group on a list.

PhotoMesa implements two algorithms for space filling, one is Quantum Treemaps by the author, a variation of the Ordered Treemap that is aware of the constraints imposed by the use of images, like grid alignment and common size, and Bubble Maps which is designed to have the least wasted space possible.

This work brings very interesting ideas to the table, but it has a somewhat limited reach, with only referring to “over 500 images” on the application and by only taking simple metadata, like the path and modification date of images.

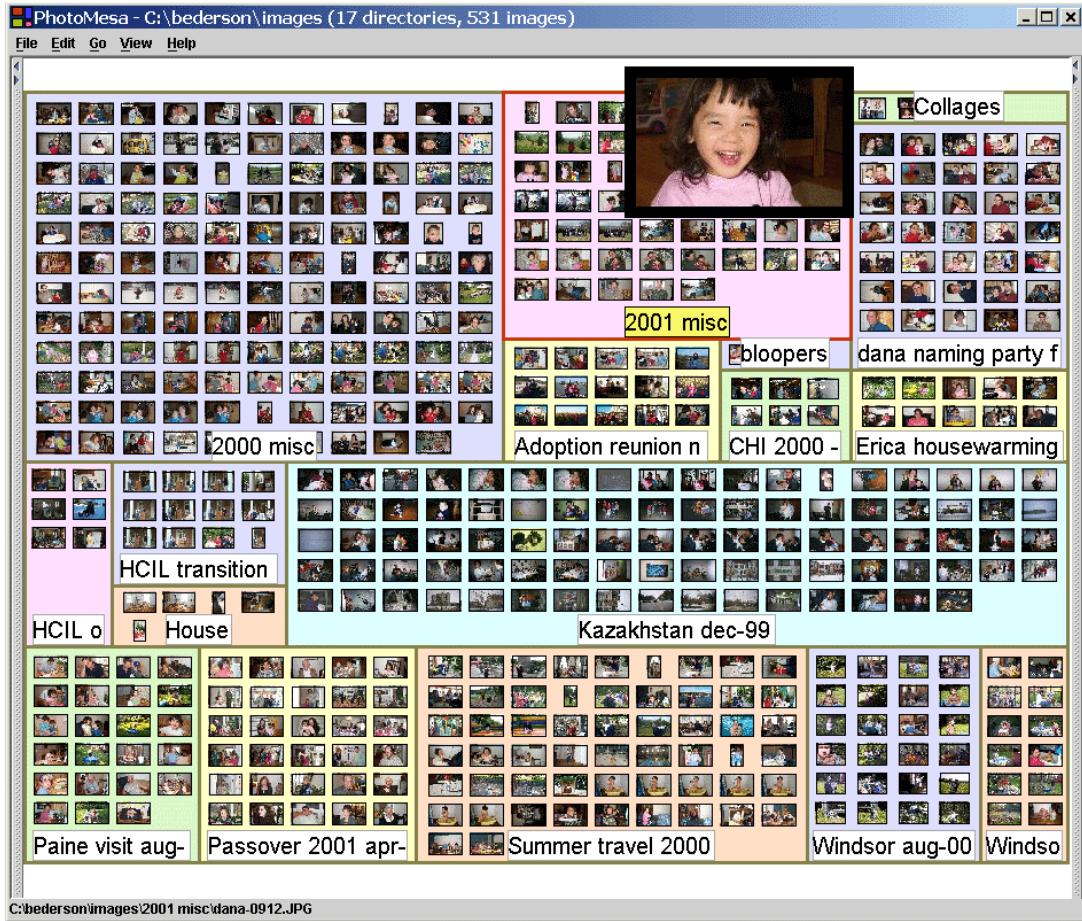


Figure 2.11: PhotoMesa with over 500 images in 17 groups.

## 2.2 Discussion

Currently there are a lot of approaches to image organization and each has its own differences as demonstrated on Table 2.1. Our survey went by multiple works and revealed various methods for extracting the features from each image, multiple visualization methods from grids to flying images.

One of the main problems is obtaining useful information from low level feature extraction of the image contents. Some try to get the most out of each image, with a variety of complex and time consuming procedures, e.g., a variety of methods based on color histograms [55, 58, 12, 54] or even more detailed ones with textures and convolution filters [30]. Others try to focus on avoiding the complex computations by only getting simple, but somewhat useful information [48]. To contrast with these methods, Grgensohn's work [27] has found that users prefer other ways to filter the collection, like tags, dates and locations. Low level features are still used, but they have to be kept to understandable and useful options.

Date and location are simple similarity measures and can be used to group the collections by events and locations like Naaman did on this work [43]. Current mainstream software like Apple's iPhoto<sup>1</sup> or Google's Picasa<sup>2</sup> are already doing it in a semi-automatic way.

<sup>1</sup><http://www.apple.com/iphoto>

<sup>2</sup><http://picasa.google.com>

Work	Organization				Visualization	Focus	Size
	features	date	gps	tags			
1. Qiu [48]	simple color measures	—	—	—	Grid	Simplicity and Efficiency	60,000
2. Rodden [51]	—	—	—	✓	Grid	Similarity usefulness	100
3. Porta [47]	—	—	—	—	Spot (and others)	Unconventional visualizations	400
4. Cooper [16]	—	—	—	—	Static and animated	Usefulness of animations	—
5. Strong [55]	color histograms and gradients	—	—	—	SOM	Evaluation of different features	2,200
6. Tan [58]	color histograms of subdivisions	—	—	—	—	Evaluation of different features	12,000
7. Naaman [43]	—	✓	✓	—	—	Organization based on events	N/A
8. Chen [12]	colors, edges and textures	—	—	—	—	Efficient fast-sparce clustering	10,000
9. Heesch [30]	six different features	—	—	✓	Radial	Complex similarity network	32,000
10. Phorigami, Hsu [34]	—	—	—	—	Grid with groups	Interaction on gruppded photos	1,333
11. Schaefer [54]	color histograms of subdivisions	—	—	—	3D Sphere	3D mapping and interaction	4,500
12. MediaGLOW, Grgensohn [27]	—	✓	✓	✓	Overlapped graph	Having the best way to find photos	450
13. PhotoMesa, Bederson [3]	—	✓	—	✓	Q. Treemap / Bubblemap	Zoomable browser	500

Table 2.1: Different browsing methods

Textual metadata like tags and descriptions are also widely used both on our survey [51, 27, 30, 3] and on all mainstream software. The problem with tags and descriptions is that people usually don't assign them to their photos, but that's not a problem we're interested here.

The visualization is the field with most variations and experimentations, like Porta's work [47] where multiple options were tested, but only a couple of the most simple were considered useful. The 3D Sphere [54] is also visually interesting, but doesn't provide a better interface to the collection since it's based on a grid view, but hiding many images on the far side of the sphere. The Phorigami work [34] introduces some interesting metaphors for reducing the space occupied by some groups of photos, but some make the visualization harder and could, therefore, be improved.

Grgensohn and Bederson's works are among the most interesting and relevant to our vision. MediaGLOW's [27] visualization approach allows images to be organized by various features and to be filtered down, displaying matched photos alongside placeholders for photos that are only match partially. The

timeline on the top is also very useful. It has some problems like image overlap and capacity for showing large collections, but the ideas are still interesting. PhotoMesa [3] has a visualization style close to what we want from our work, but its reach is short, i.e., has a small set of features and dispositions, doesn't handle many images and the UI seems a little cluttered with all the strong colors and borders.

The seemingly more useful exploration tools are the ones that handle less images, which seems like a contradiction. The more images a system can hold and efficiently display, the more useful it gets. From our user survey (appendix A), we learnt two thirds of people have less than 10 000 photos (A.2) and we will do our best to reach that level.



# Chapter 3

## Solution Requirements

Our survey was based on various types of previous work from the last ten years. Image browsers and technology have evolved a lot since then, but there still isn't a definitive way for a user to look at his larger photo collection and understand its content and evolution.

We have the vision of a system that displays a large set of user's photos at the same time, in various arrangements, revealing patterns, differences and similarities between them.

We will now expose this vision by presenting the main goals we want to achieve with this thesis, as well as some of the guiding implementation requirements that we followed for a better end result.

### 3.1 Main Goals

The main goal of this thesis is to provide a different approach to the photo collection browsing methods. More specifically, the work should:

#### **Extract interesting information about the images**

We want to be able to organize and classify images. As we have seen on the Related Work, there are many ways to extract information from images, some simple, some complex. We have chosen a few that we think are the most relevant to the scope of our work. We will explain what we chose and why below, in section 3.3.

#### **Provide an interaction with the full set of images**

We want to enable the users to easily view and interact with the all the images at the same time, if they want to, instead of just limiting to a small window of images. Many of the Related Work did this and displayed all images at the same time [48, 12, 28, 3]. Others only displayed a subset of images, but we think the users should have the power to create and view their own subsets from the full one.

### **Efficiently display a large number of images in a single screen**

This goal goes inline with the previous one: to be able to interact with a large set of images, we should be able to display them all, at the same time, using techniques to make better use of available screen space. We saw different ways to display images on the Related Work, some focused on displaying images efficiently [48, 51, 43, 34, 3], others more focused on the relations between the images [28, 30]. Hsu's work [34] is specially interesting for its idea of merging images that were related and we want to include this in our work, but with an automatic way of grouping and a clearer UI.

### **Allow the manipulation of the display**

We want the users to create different views that enable new perceptions of their collection, based on the extracted information. Many works allowed some sort of highlight [28], zoom [3] or similarity filter [48, 30]. Some are more useful than others and we think we should apply them sensibly. Since we are going to display a large amount of images at a small size, there will be the need to increase the size of the images, therefore we require the ability to zoom in on images and to provide some filters for any available data, but also for custom user selections. It is also interesting to have some similarity filtering.

With this capabilities, our work should be able to provide the user not only a better understanding of the collection, but also with an easy and interesting way to visually combine and view photographs.

## **3.2 Implementation Requirements**

In addition to the referred main goals, we set ourselves some design goals, or requirements, for our implementation. With them, we want our work to get closer to a real application, that real users can use and have some flexibility for it to evolve with time.

Therefore, we set the following requirements:

### **Ease of Use**

Ease of use is one of the most important characteristics of any piece of software and can shape how well it will sell. Much more attention has been given to the user experience (UX) in the last few years. Systems that are easier to use, get the job done faster, are more enjoyable to use and allow less experienced users to use them.

Since this is such an important characteristic, we aimed at providing a simple interaction from the beginning to the end, while also providing a powerful system, even though it could be even more refined in a few areas.

## **Extensibility**

The extensibility factor of a system is also important for the added value that can be obtained by quickly adding new features, either by the developers or by third parties. We made some parts of our work with this in mind, by allowing either external plugins or by generalization of code, allowing for future improvements with less trouble.

## **Performance**

One of our main goals is to display a large set of images on the screen, but this brings problems since each image, in full-size, can take a good set of resources of the system. If we multiply this resources for a thousand images, we will not have a performance work.

Therefore, we set this requirement for having a performance work, that can be used with at least a few thousand images without taking down the system while using it.

## **Persistency**

Our system will spend sometime generating and gathering data for each of the available feature extractors, for each image. To avoid having to re-do work in case of a failure, addition of more images to the collection or stopping to resume later, the system must store the data after its generation. This should be accomplished using a system wide framework to provide easy storing of all the data. Since this work is intended to be an exploration of visualization concepts, so we will not support interoperability from our system to others by using MPEG-7 or other metadata descriptor standards, but instead use something that is simpler to implement.

## **3.3 Features of Images and Photographs**

As referred before, we want our system to extract interesting information about the images and we will now explain this in detail.

Unlike textual data, images are a type of data where it is not trivial to extract information from, in a computational environment. It's easy for us, common people, to understand what certain picture is showing. We can easily distinguish if there are, for instance, animals, people, flowers or buildings, but it's hard for a computer to do the same, and it's even harder to understand if that certain photo of a building and a person was taken because of the building, the person or both.

There have been various developments in the feature extraction front [39, 21, 53] and it is currently possible to perform various detections in images with various levels of satisfaction.

Some examples of working solutions for feature detection:

## **Face recognition**

Detection of the presence of faces in the images, their position and relations between them as some other works refer [61, 13, 57, 33]. Face recognition has been gaining public appreciation in the past few years since many photo applications and services have been including it. Our user survey reflects that people like this feature and request it (A.2). The problem with face recognition is that it requires users to identify who are the persons on each photo which, generally, is a long and repetitive process. For concept, we will only detect the presence of faces and not spend time identifying them.

## **Object recognition**

Identification objects in images. This is a useful extractor, but it requires a huge library of sample imagery for comparison [60] and it's not feasible to adopt in our work.

## **Identification of perceptive colors**

What are the main colors that users perceive in certain image [56, 58]. An image is composed of different colors or tones and people can identify images by their main colors and dispositions. There are simple ways, like calculating the mean color for an image, but that's not faithful. Detecting perceptive colors identifies the different colors of an image, providing more realistic data.

## **Image sequences**

Identification of images of the same scene that were taken in sequence, for grouping purposes [17]. Images taken in sequence of the same subject should be gathered into a group of images, with the purpose of reducing the needed space to display them. This is important since almost half of the people we inquired on our survey claims to use techniques that capture a great number of photos sequentially (A.2). This can be either by looking at the capture date or by analyzing the visual similarities of sequential images.

## **Image similarity**

Identification of similar images across the whole collection enabling a search of similar images from a selected one or just grouping similar images according to defined parameters. There are many ways of doing similarity, some simple, like measuring the average color of the image [55, 48, 54], some complex like measuring various color attributes, edges, textures, applying filters or using convolution like Heesch [30] and Chen [12] did. For our work, we will explore options for color similarity, paving the way for other more accurate methods.

We want to use some of this methods to automatically extract information from the contents of images but, since this work is mostly focused on photography, there is another way to obtain really useful information that is EXIF Metadata. None of the related works seemed to use it, although users are interested in them, as can be seen on our user survey (A.2).

EXIF Metadata is a standard format for metadata included in digitally captured media, like photographs captured by digital cameras. These devices save a lot of information on the image file, like the date and time of the capture, camera settings, camera orientation, location data and even detected faces<sup>1</sup>. An example of the metadata fields can be seen on Table 3.1. This is textual data and, therefore, much faster and easier to retrieve and analyze than the content based methods we discussed above, being a great addition to them.

Mixing all this different data, users should be able to interact, filter, sort and organize their collection in various ways, obtaining new and different perspectives of their images.

Field Name	Example Value
Camera:	Canon PowerShot S5 IS
Exposure:	0.017 sec (1/60)
Aperture:	f/3.2
Focal Length:	6 mm
ISO Speed:	80
Exposure Bias:	-2/3 EV
Flash:	Off, Did not fire
Software:	Adobe Photoshop Lightroom 3.4
Max Aperture Value:	2.7
Subject Distance:	0.3 m
Metering Mode:	Multi-segment
Sensing Method:	One-chip color area
Custom Rendered:	Normal
Exposure Mode:	Manual
White Balance:	Auto
Keywords:	“Beach”, “Sand”, “Castles”
Date and Time:	2010:09:18 11:52:59
Lens:	6.0-72.0 mm
Approximated Focus Distance:	0.3
Flash Mode:	Off
Metadata Date:	2011:06:14 20:03:04+01:00
GPS Latitude:	38 deg 52' 47.67" N
GPS Longitude:	9 deg 9' 47.13" W
GPS Altitude Ref:	Above Sea Level
GPS Altitude:	214.9 m
GPS Date Time:	2010:09:18 10:52:59

Table 3.1: Example of some relevant EXIF fields imprinted by the camera and by the computer software used for post processing.

---

<sup>1</sup>the availability of some of this data requires capable camera hardware and software that are getting more common nowadays.



# **Chapter 4**

## **Eagle Eye**

Eagle Eye is the product of this thesis and is our implementation of the requirements previously detailed. It is a visualization tool that enables the display and manipulation of a large image set at once. It is focused at the regular computer user that has a few thousand digital photographs stored on the computer. It allows navigation, through panning and zooming of the canvas where the image collection is disposed, sorting in different ways while also allowing filtering, either textual or visual.

In this chapter, we will go through how does Eagle Eye work, what are its components and how they interact.

### **4.1 Overview**

Our requirements stated the need for features to be extracted from images, so we started by creating a system to enable that. We then changed our focus to the visualization part of the work and redefined the general architecture of the work. The system is now composed of two main elements: the Backend, which is the system that extracts features from the images and prepares everything to be displayed; and the Visualization which performs the display from the images and their extracted features to the user.

This is the overview of the system that we will now detail in the following sections.

### **4.2 Backend**

The backend is one of the two parts that make Eagle Eye. Its purpose is to extract information from images and set everything up for the Visualization.

Currently it is a command line utility that allows the user to enter paths for folders containing image files. The system will then read those images, gather their metadata, process them with the existing plugins to extract visual features and, finally, generate and output the multi-scale imagery and the control metadata required by the visualization.

We will now detail its architecture, and implemented feature extractors.

#### 4.2.1 Architecture of the Backend

The Backend comprises a library manager to hold the images, feature extractors to process those images and persistence to save all generated data.

Figure 4.1 is a simple explanation of the components. The Eagle Eye part is the main application, containing the library and feature extractor managers. Both deal with files on the disk, JPEG image files and DLL extractor plugins, respectively. The user interacts with the core of Eagle Eye which currently provides a command line interface for its actions, like the image import and plugin execution. The import gathers the files and their metadata, to be later accessed by the plugins for processing. Plugins store the resulting data inside the library manager and can be accessed afterwards for outputting by a special plugin.

We will now go through this components in more detail.

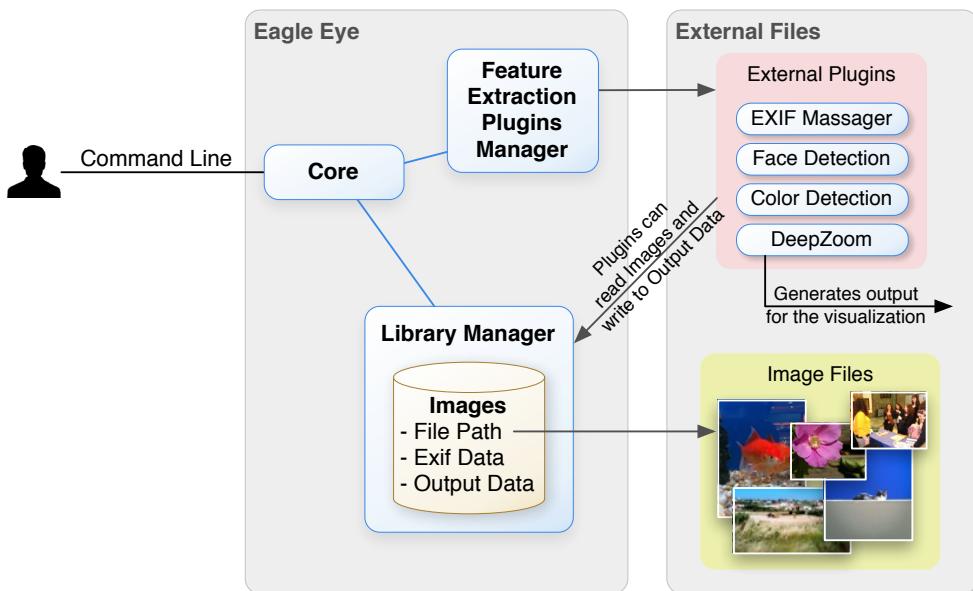


Figure 4.1: Basic architecture of Eagle Eye's backend

#### Library Manager

The system displays images and, therefore, it needs to know what to show. This is where the library manager comes in. It creates a database which indexes existing JPEG image files stored on the user's computer and makes this information available for other modules to use. It is designed to be used with digital photographs which contain the aforementioned EXIF metadata like time, date, camera information, or location. This information is gathered upon import and is available throughout the Backend.

We explored a few ways to develop the image import process and we rested at the fastest we found. The user refers a folder to be imported and we use a third-party program, ExifTool<sup>1</sup>, to crawl through the user-specified folders while identifying all the JPEG images and returning their Exchangeable image file format (EXIF) metadata which is then stored by our system (fig. 4.2).

<sup>1</sup>ExifTool is a utility that allows easy read and write of file metadata. <http://www.sno.phy.queensu.ca/~phil/exiftool>

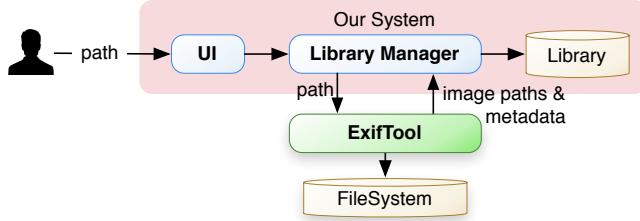


Figure 4.2: Process in use by our system for importing a folder using ExifTool.

Another option for importing metadata was to invoke ExifTool as part of a feature extraction plugin (FEP). Although that could fix a couple of problems with the current implementation, it doesn't make the metadata as ubiquitous as needed. Most FEPs rely on some EXIF plugins to work correctly and the current implementation doesn't easily allow inter-FEP data-sharing.

### Feature Extraction

To enable the Visualization to arrange the images on the screen in different ways, they need to be classified. Some information is easy to obtain and compare, like when the photograph was taken. Other information needs to be extracted, like the number of people in the photo or what are the most relevant colors in the image. Fig. 4.3 is a short example of what feature extraction is all about.

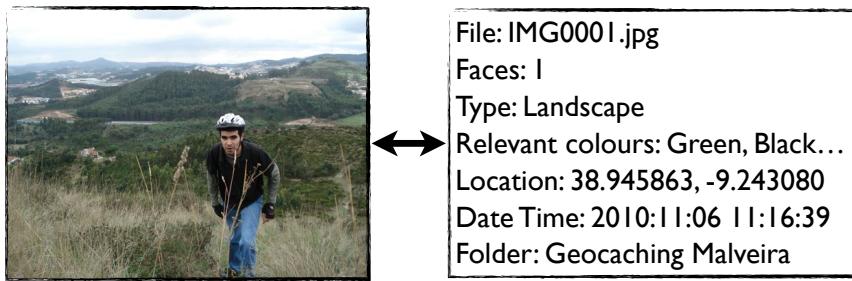


Figure 4.3: Example of information extracted from an image.

We had a few ideas for extracting features from images and, to be possible to add more along the way, we developed a plugin system to ease the creation of other feature extractors in the future.

Each Feature Extraction Plugin is separated from the main system. They need to implement a common interface and are given the ability to access the image data from the library, and save the computed data back. They have freedom to access the image file or its EXIF information. They should, in the end, store the processed data in a specified way so it gets exported to the visualization. Implemented plugins will be explained in section 4.2.2.

### Persistence

Persistence of both library and plugin data are required so the system doesn't lose information that took time to generate. To ease the interaction with a database system, we created a database abstraction layer that hides the complexities of interacting with said system. It also allows us to change to another

database system if we see the need for it. We chose Oracle's Berkeley DB<sup>2</sup> for its speed in retrieving data.

With this layer we can hide some optimization complexities like lazy-saving and lazy-loading. We use lazy-saving to save data to disk by chunks instead of doing it on each small update, speeding up the update process. Lazy-loading is not yet implemented, but it is essential with libraries with tens of thousands of images, where keeping a complete library in memory is not feasible.

### 4.2.2 Feature Extraction Plugins

Back in the Solution Requirements chapter (3.3), we detailed some possible features that were interesting to be used in this work. In this section we will detail what extractors have been implemented.

Currently, we have four feature extraction plugins:

- Selection of useful image metadata
- Detection of image's main color
- Face detection
- Generation of multi-scale imagery

We now proceed to the explanation of each one of this plugins.

#### Selection of useful image metadata

This plugin acts as a filter for all the available EXIF tags. It picks the most relevant ones, codes them in a pre-defined way and appends them to the rest of the information to be exported for the visualization.

Information when the photo was captured, what device was used, the path where it resides or information about the location where the photo was taken are a good examples of the most commonly relevant tags, as can be seen on the user survey (A.2). In the future, this set of extracted tags could be optionally set by the user.

#### Detection of image's main color

Sometimes people don't recall where or when a photo was taken, or where is it, but they vaguely remember that the photo had some dominant colors, like the red of a parrot in a green background (fig. 4.4). This kind of information can be helpful when searching for a photo in a collection.

Color extraction from images has been a long standing problem [62, 55, 26, 28, 59, 20, 10]. We wanted to extract the most perceptible colors from images so, in the parrot example (fig. 4.4), the system should associate the image with the colors red, green, white and black and avoid colors that have little relevance, like the small blue of the feathers and also ignore small tonal variances in the reds and greens.



Figure 4.4: A colorful photo.

<sup>2</sup>Berkeley DB is a high-performance, embeddable, key-value, file-based database available at <http://www.oracle.com/technetwork/database/berkeleydb>

For that we explored two methods that reduce the number of colors in an image to the most essential ones, the first being an **adaptive method**, selecting a few averaged colors from the image and the second using a **predefined color palette** as reference to select the colors.

A common JPEG photo can contain 16.8 million different colors. Our adaptive method reduces the possible colors to less than ten<sup>3</sup>. The obtained colors are chosen by averaging the colors in the image, so if there's a lot of blue tones, a single, averaged blue will be replacing those tones (see the middle image of fig. 4.5). In areas that have little tonal variation of the color, the resulting color will be very similar to the original (like in the right image of fig. 4.5). But sometimes this method fails to save important colors when they occupy a relatively small area of the image or if the image doesn't have a strong contrasts (shown by the left image in fig. 4.5).



Figure 4.5: Three images where the left half is the original version and the right half is a version with very limited number of colors provided by the adaptive method.

The second method uses a predefined color palette (fig. 4.6) which is an adapted mix between the eleven most recognizable colors (red, yellow, green, blue, purple, brown, orange, pink, black, white and grey)<sup>4</sup> and the 21 ColorAdd colors<sup>5</sup>. This mix has a great range of colors and each one can be easily named with recognizable words<sup>6</sup> which then could be assigned to the images. The images processed with this method get their colors changed to the most similar ones present in our palette. This assures that colors in a relatively small area or images with low contrast still get the deserved attention (fig. 4.7.3), unlike the adaptive method (fig. 4.7.2). This method doesn't reproduce the colors so well as the adaptive but, since we want to obtain generic colors, this one is more reliable.



Figure 4.6: Color palette in use for restricting the possible colors in an image.

This was the research we made to extract perceptible colors from images. We generated the above

<sup>3</sup>We tried with multiple options from one color to ten colors and the results vary from image to image. We think eight is a good value, keeping a nice set of colors on colorful images.

<sup>4</sup>Qiu [48] explains these are the most universal recognizable colors in any language.

<sup>5</sup>ColorAdd is a 21 color catalog with symbols for each color designed specifically for color blind people. We used the colors as a reference, specially the light and dark variations.

<sup>6</sup>We want to avoid names like "Amaranth" or "Munsell" that most people don't know about. A large list of these names can be found on [http://en.wikipedia.org/wiki/List\\_of\\_colors](http://en.wikipedia.org/wiki/List_of_colors).



4.7.1: Original image

4.7.2: Adaptive method

4.7.3: Mapping method

Figure 4.7: Comparison between the original image, the adaptive color reduction method failing to keep the pink and the blue colors and the color mapping method.

demonstration images using ImageMagick<sup>7</sup> and its “colors” feature for the adaptive method and the “remap” for the palette method. Unfortunately, time was not on our side and we had to move to a simpler system.

The current color extractor plugin does a simpler job of calculating the median color of the image and use the hue to index images.

We took some ideas from the work of Qiu et al[48] of exploring a method of image indexing that is fast and simple. On this work, images are sorted into “bins” according to their average color. These bins are divisions of color planes, indexed using binary trees allowing for very fast searches.

We used the open source library AForge.Imaging<sup>8</sup> to obtain histograms for the images and compute their average color in RGB<sup>9</sup> and HSL<sup>10</sup> values. Both those values are then stored on the plugin-extracted data for each image. The images will be distributed into bins on the Visualization. We could assign bins at this point but, by not doing so, we are giving freedom to the users to change the number of bins used for display. We will discuss this on the Visualization sections ahead.

Although we didn’t implement this plugin the way we desired, it demonstrates that color extraction is feasible and, with more time, a better method that does more than just averaging the colors, including detecting the most relevant ones, could be implemented.

## Face Detection

The face detection plugin is based on the open-source OpenCV library<sup>11</sup> which processes every image file and detects existing faces (fig. 4.8).

No face recognition software is perfect. Usually if the software can detect every face, it will probably detect some other things in images that aren’t faces (false positives). If it’s successful in only detecting

<sup>7</sup>ImageMagick is a software suite to create, edit, compose, or convert bitmap images. <http://www.imagemagick.org>

<sup>8</sup>AForge is a framework designed for developers and researchers in the fields of Computer Vision and Artificial Intelligence <http://www.aforgenet.com/framework/>

<sup>9</sup>Color model composed by red, green and blue <http://en.wikipedia.org/wiki/RGB>

<sup>10</sup>Color model composed of hue, saturation and luminance [http://en.wikipedia.org/wiki/HSL\\_and\\_HSV](http://en.wikipedia.org/wiki/HSL_and_HSV)

<sup>11</sup>OpenCV (Open Source Computer Vision) is a library of programming functions for real time computer vision. Available at <http://opencv.willowgarage.com>

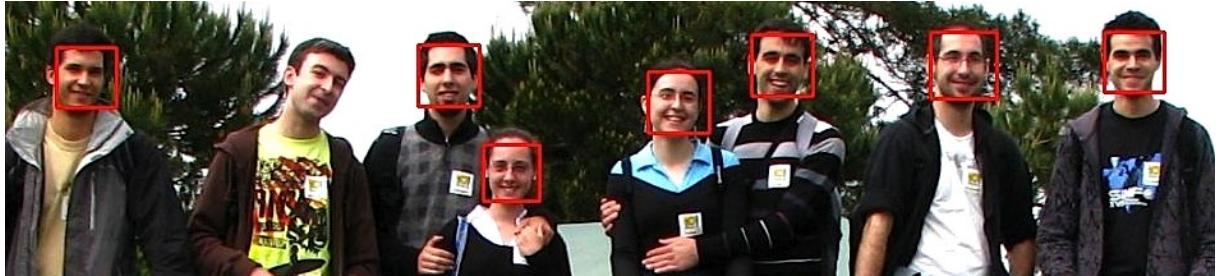


Figure 4.8: Example of the OpenCV library detecting faces on a common photo.

faces, it will probably miss some other faces that aren't ideally positioned (false negatives). OpenCV is included in the latter, only detecting faces, but also missing some that are tilted (like in fig. 4.8) or turned on the side.

This process is quite computationally expensive and therefore we resize all the images down to a more acceptable size, making the process more than five times faster.

We tested 29 images, from six different cameras, ranging from one to ten megapixels, and containing up to thirteen faces. The test consisted in running face detection on each image, in its original size and in various resolutions from 2000 to 200 pixels on its longer side, comparing the number of faces recognised and the time needed to process them. The results can be seen in fig. 4.9.

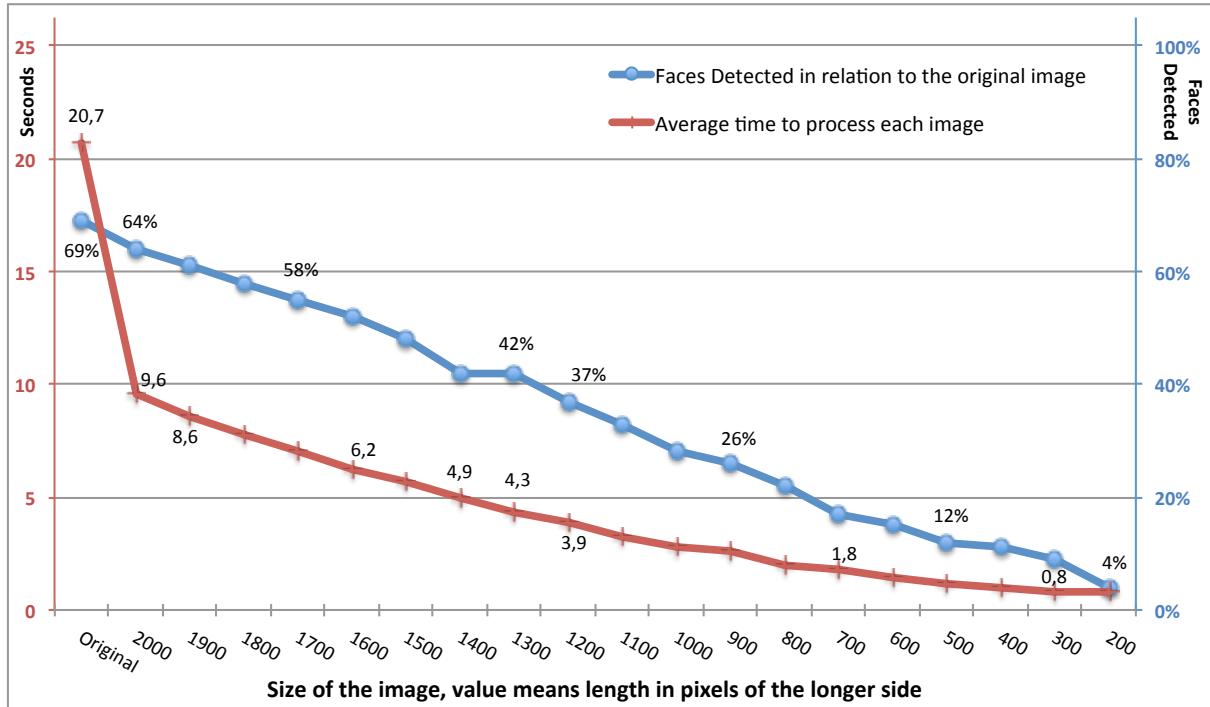


Figure 4.9: Detection vs. speed results of the face detection test. 100% of faces is the actual faces present in the test images.

The purpose of this test is to identify how much we can reduce the images while maintaining a high recognition rate, we are comparing the recognition results of the downscaled versions to the original size and analyze the speedups and failures in recognition. We do include the number of faces actually present in the photos for comparison, corresponding to the 100% value in fig. 4.9.

We can see that by only reducing the images to 2000 pixels on the longer side, the processing time fell to less than half (20.7 to 9.6) without much loss in recognition (69% to 64%). The 1300 pixels was the chosen value for being the last with more than 40% recognition rate (60% of the full size image) and being 4.8 times faster<sup>12</sup> than using the original image. In the future, with more tests, we can fine tune the resizing algorithm to get better results.

### **Generation of multi-scale imagery**

This plugin generates all the data files needed for the visualization to work. The visualization relies on the DeepZoom technology and it needs to process the images before they can be displayed. This plugin does exactly that.

Using a library from Microsoft, the plugin generates, for each image, a metadata file and a set of image files representing the original one at multiple scales, from a single pixel to a large, detailed image.

After passing through all images, the collection as a whole is subject of additional computation, this time generating imagery for all images as a single set and a metadata file that agglomerates all image sets used. This metadata file for the collection (called collection.xml) is then altered by the plugin to attach to each image, the data previously generated by the other plugins.

## **4.3 Visualization**

The greatest challenge of this work was the creation of the visualization part for its requirements. We will now describe this part of Eagle Eye, its architecture, visualization techniques, sorting and filtering capabilities.

### **4.3.1 Overview**

We wanted to make the visualization simple and easy to use, while keeping it flexible and capable enough to allow for an enjoyable and relevant experience.

After the backend has finished the all the processing that is needed, the visualization can be opened and all images that were added to the backend's processing list will appear in what we call "the canvas". After loading the metadata, the user is also presented with a set of options on the top toolbar, which contains all the controls needed to use the system (fig. 4.10).

#### **The Canvas**

The canvas is the most relevant part of the visualization as it displays the user's images.

According to A. Torralba [60], an image with 32 pixels per side is the minimum size that allows a user to recognize what it is. His work was about an image library that had nothing to do with the user's photo. On Eagle Eye, we can have even smaller images because our focus is the user's own library. This allows much easier recognition of images since, by looking at the form and colors, the user recalls

---

<sup>12</sup>4.3 seconds per image versus 20.7; one hour and 10 minutes per thousand images versus almost six hours



Figure 4.10: Example of the visualization of 5679 images, organized by capture date, in Eagle Eye.

what is that image without having to “read” and understand a totally new image. Another very important help is the grouping of events (e.g. by date or path) which allows the user to extrapolate the contents of the whole group from looking at a few of them instead of trying to understand each image individually. This allows much faster recognition of groups with images at really small sizes, even if not all images can be understood.

An example of this can be seen on figure 4.10, which is a screenshot of the application running on a 13.3” screen with a resolution of 1280x800 pixels, where the larger side of each image uses 15 pixels and takes 3 millimeters. On this example, different groups are clearly distinguishable from each other and most of them are easily recognizable by the owner of the photographs. Each group seems to have a specific set of colors, e.g., blue, green, yellow, orange, gray or black, making easy for the owner to distinguish and recognize the photos.

A photo application that only displays all images isn’t very useful, therefore, manipulation of the canvas is allowed by zooming and panning. This means the user can, at any time, use the mouse to drag the canvas around or, by clicking or scrolling, zoom in and out of the canvas. Zooming goes from the default view of thousands of images at the same time, to the full screen view of one of them, and everything in between in a smooth way. This allows a closer look to any group or image, providing an easy and fluid way to navigate the collection.



Figure 4.11: The toolbar on top of the visualization UI.

## The toolbar

The user can use the functions on the toolbar to filter and sort differently. The toolbar is divided in three sections: History, Display and Filtering (fig. 4.11).

The **History** section contains some basic functions that work similarly to the current web browsers. There are buttons for going back and forward between the history of display states and a save button for bookmarking the current display state, allowing the user to easily get back to it later.

The **Display** section, in the middle contains two options to change the image display: the sorting options and the display overlays button. The former presents the available sorting options for the current collection, based on the available metadata and on the best ways to display them. The content is presented according to the selected option and changing to another causes the images on display to move around to a new position and form a different display order. This sorting and disposition options will be explained in a later section. The other button in this section of the toolbar is called “Group Names” and enables or disables a layer of information on top of the images. This layer distinguishes the groups of images in display by painting them with different colors and presenting names for them, depending on the selected sort option (fig. 4.12 shows an example of this). Grouping will also be explained bellow.



Figure 4.12: By enabling overlays, the user can see names for the groups and a much more clear group distinction.

The third and final section of the toolbar is the **Filter** section. It contains controls to filter images by using simple text and to visually select images on the canvas. This options will also be explained bellow.

### 4.3.2 Disposition of Images on Canvas

The different ways to dispose the images on the canvas was a matter that required some exploration of possibilities, like the ones seen on the Related Work ([3, 7, 12, 27, 30, 34, 47, 51, 54, 55]).

We have come up with a some options for photo grouping, sorting and filtering. We kept a sensible set of the choices so it's easy and understandable for users. To make sense of how all this pieces work together, we are going to explain how do we sort images into groups and the different ways the user can arrange those groups on the canvas, followed by how filtering is done.

#### Sorting images into groups

As we've seen previously, the backend outputs metadata for each image. This set of metadata is loaded into the visualization application of Eagle Eye and is then indexed by type.

For instance, each image has an associated creation timestamp which will be aggregated by days, generating an image group for each day.

Similarly, the mean color associated with each image is indexed and groups are generated by splitting the hue spectrum. We used a quick approach and divided it into 12 equal bins with names, as can be seen in figure 4.13. From the image, it's possible to see that some colors have a greater usage of the spectrum, like green, which takes the most part of three bins, while other colors use only one. This gives the user multiple tones for some colors and only one for others. We made it this way so the number of bins could be quickly changed to something else if we saw fit, but the most sensible option might have been to manually define the bins with different sizes and more accurate names.

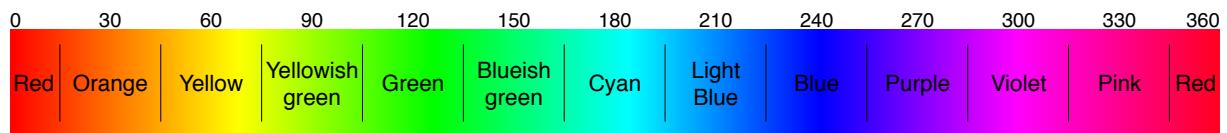


Figure 4.13: Division of the hue values in 12 equal parts, 30 points apart from each other.

Another option is the grouping by device name, which usually allows to distinguish between who took the pictures if, for instance, different people have different cameras on the same event.

The last option currently available to the user is grouping by path which groups together the images that were already grouped by the user, on the file system. This allows for the display of an organization that is recognizable to the user, which can make a good starting point.

On the canvas, group boundaries are identifiable by discrete gray borders and, when the Show Overlays function is active, by color rectangles that also contain the groups' names.

#### Different dispositions

After having the images grouped by any of the sorting options referred on the previous section, the system has to know how to display them on the canvas.

We looked into various options ([3, 7, 12, 27, 30, 34, 47, 51, 54, 55]) and picked the concepts that we thought that made sense and that would be easier for the user to understand. We chose a

treemap view and a column-based linear view. Both of these are grid-based layouts, meaning that images are positioned inside a defined grid on the canvas. We also looked into free positioning systems ([27, 51, 54, 55]), but they make it harder for the user to understand the images within, since some images will be covered by others. When displaying thousands of images, it's important to make them easy to see, and mixing them up wasn't the appropriate thing to do. We focused on other ways for making it easier to the user see what matters and we will talk about them later on.

The first layout technique we employed was the treemap. Introduced by Brian Johnson and Ben Shneiderman [36] in 1991, treemaps recursively subdivide areas into rectangles. The problem is that areas can take many forms, from squares or large rectangles to almost thin lines, like can be seen on fig. 4.14. Applying treemaps to images calls for the adaptation of the algorithms to make sure the areas can correctly hold the images and that all images in all groups have the same size and are positioned in the same grid, to make them easier to view. This ideas are supported by B. Bederson on his work with the Quantum Treemaps [3]. We tried to apply quantum treemap as our treemap algorithm but, due to its complexity and other recurring problems, we had to move to a simpler method and adapted the squarified treemaps [7] by Mark Bruls to support the grid of images that we required. This algorithm was much easier to understand, adapt and implement.

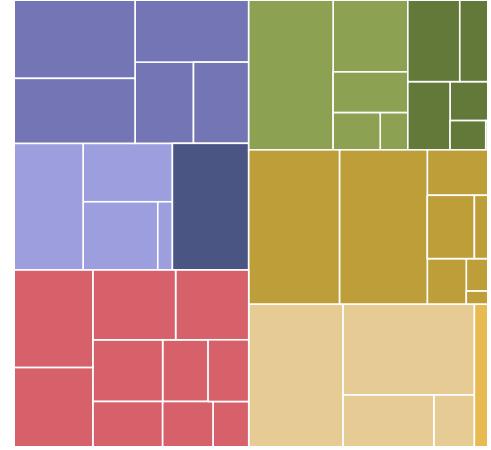
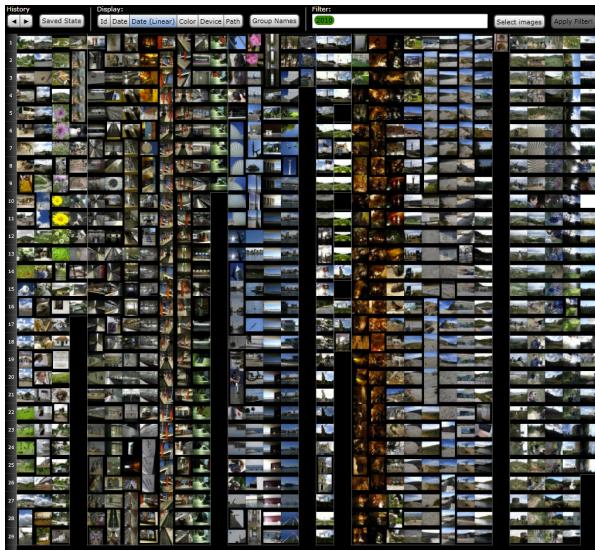


Figure 4.14: Common treemap based on space filling.

Our resulting treemap algorithm displays larger groups first, leaving the smaller ones to the end, making an effort to layout groups as rectangles with an aspect ratio as close as possible to the screen's aspect ratio, for filling the screen when the user zooms on one. There's also an effort to fill the space left between larger groups and the edges of the screen, making the display more compact and with less empty space across the screen edges.

One problem of the squarified treemap display is that group sequence is irrelevant. Groups are positioned by their size, which is unacceptable for sorting options that require some sequence, like sorting by time. For this we created a linear display (fig. 5.1) that uses columns and displays groups sequentially. Each group may fill part of a column or various columns, depending on their size. With the aim of reducing wasted space, groups that fit on the wasted space left by the previous group use that space to display themselves. This is useful to group the few pictures the users may capture of their day-to-day, between days (like weekends) that he went on a trip and took a much larger number of photographs.

Currently we are using this layout system only for the date display since the use of columns makes visualization harder, requiring either some panning from top to bottom or selecting the group using the filter tools explained ahead. This is an area we must improve, and we will discuss some ideas later on.



4.15.1: Images only.



4.15.2: With group divisions overlaid.

Figure 4.15: Example of the linear display.

## Filtering

In addition to the sorting options explained above, Eagle Eye also provides the user with the ability to filter images. This allows the users to focus on specific groups of pictures that are of interest at a particular moment.

For instance, the user might want to only see photos of a day, or a person, or person in a specific day or even pictures that are mostly blue.

For all this, Eagle Eye provides two ways to specify this kinds of constraints: using the filter bar or selecting pictures on the canvas, and we will now look into both them.

The filter bar's purpose mimics a regular search box, similar to the ones that exist on applications like internet browsers, file browsers or photo browsers: it provides a way for the user to type what he is looking for and get some suggestions to help him with the search. Since we wanted a simple UI, we figured that we had to provide smart suggestions to the user, so that he can feel comfortable using it for multiple types of searches. Searches are performed immediately upon selecting the desired filter. Currently, searches are only intersections (commonly “and”) of the entries in the filter bar, but after adding adequate UI it will be possible to also use unions (commonly “or”).

Every piece of the metadata can be used as a filter and the suggestions list shows what are the available options and what type of metadata they relate to, for instance, when looking for a person's name, the “keyword” metadata type is shown, when searching for a place, both “keyword” and “path” types might be presented, if the path to the pictures included the place where they were taken.

This examples are quite simple, but we wanted to go further and allow searches like “Summer”, for all photos taken around the summer time in all years, “May 2010”, for all pictures taken during that month, “Has 2 people” for photos that have been detected to have people in them. Although some of these examples have not been fully implemented yet, they are in the plans. See figure 4.16 for some examples.

The second filter method allows the manual selection of images from the canvas, either by a common mouse drag-and-drop on the canvas or by selecting entire groups with a mouse click. This two interaction options can be activated on the respective button on the end of the toolbar.

As explained above, the filter bar intersects its elements immediately after being added, but that behavior in the function would make it impossible for the user to make multiple selections, therefore they are only unified and the filtering will only be performed when the user presses the Apply Filter.

Since the selections are displayed in the same way text filters are, the removal of a selection is made in the same way, through the filter bar.

### 4.3.3 Architecture of the Visualization

We will now take a look into the architecture of the visualization part of Eagle Eye.

#### Metadata indexing

At the center of the visualization is the DeepZoom canvas. By default, it displays images in the order they are specified on the “collection.xml” file, the file generated by the backend that identifies the multi-resolution imagery.

The “collection.xml” also contains some metadata for each image, added by the backend. The visualization parses the file and extracts this metadata which is then processed and indexed.

The backend encodes textual key/value pairs of metadata in JSON for each image and the visualization decodes and aggregates them on various indexers, one indexer for each type of metadata. Eagle Eye has a few indexers that know how to read and process different types of data:

- a generic indexer, for basic text or number values
- a date indexer, that gathers images by day
- a color indexer that gathers images by color
- a keyword indexer that indexes images by each one of the associated keywords
- a path indexer that indexes images by the folder where they resided when added

The indexers are managed by a metadata collection manager that handles, for instance, the XML parsing and creation, management and retrieval of indexers.

The indexers are the base for the sorting and disposition techniques we’ve covered on 4.3.2. The indexers can have disposition types associated with them, like the date indexer, which has the treemap and the column dispositions associated with it. Other indexers only have the treemap disposition (e.g.

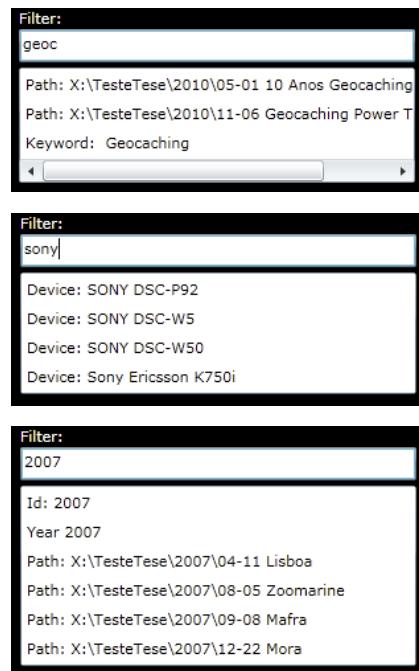


Figure 4.16: Examples of filter suggestions.

Color, Path) and others don't have dispositions at all, like the Keyword indexer that is only accessible from the Filter bar. This indexer to disposition association is our decision and aims at providing the user a simpler interface while displaying images in the best available way.

The indexers also provide data for the filter bar's suggestions to provide the suggestions we have already seen on section 4.3.2.

## Canvas

The canvas is the most visible part of this work. Its core is from the DeepZoom framework and we use it to display lots of images in certain positions, according to the state of many options and variables. To hold all those options together and to help developing some more, we have created a state manager for the canvas.

A new canvas state object is created for each different set of options for the canvas. The object stores the options and computes the positioning of all images by picking those that have been selected by the filter, get the groups for that image set from the selected indexer, dispose them using the selected disposition algorithm and finally computing and saving the position for each image.

We use memoization to avoid repeated calculations when changing display types. If a certain display has already been calculated, it will reuse the information on the saved state object to rearrange the canvas without going through the position calculations again, allowing a faster interaction.

The navigation buttons on the toolbar interact directly with the state, for instance displaying a previous state or one that was saved.

## Reducing Clutter

This work displays a great quantity of images at the same time and we thought that we could help the user by reducing some unnecessary clutter.

Sometimes users take more than one photo of the same thing and they do it for different reasons, some of them simple and others more technical. Some people just didn't like the previous photo or they want to make sure they have a few options to pick the best one and so they take more. Bursts of photographs are a similar case of this behavior, but resulting in a larger image count.

Some more advanced photographic techniques require that the user takes more than one shot of something for merging in post processing. Examples of this techniques are panoramas, where a sequence of shots are merged to create a bigger image that what the camera lens is capable of, exposure bracketing, to overcome partial under or overexposure or focus bracketing, to increase the area of the image that is correctly focused.

Various shots of the same subject are generated by all of this motives and they clutter any image browser since they end up being multiple copies the same thing.

To help the user abstract himself from this similar images, we decided to come up with a way to collect them. Hsu et al. [34] used the idea of image grouping, but probably not with a great implementation since it can be hard to see what is in the groups when they are collapsed and grouping was done manually. We chose a different path.

For our implementation, we decided to group images automatically based on the proximity of their capture timestamps, so every two images that were taken in a time span less than  $t$  seconds are grouped together. After some experiments we decided on  $t = 4$  seconds which is more than enough for sequences of images that were taken automatically, gives some margin for manually taking another photo of the same thing, but is short enough for a regular user to take a picture, look at something else move the camera and take a picture of that.

For the display we decided to show all images of the group in the space of one, but displaying the first image at regular size and reducing the others to a smaller size, perceptible when looking closer and keeping the ability to zoom in on each of them and see them in fullscreen. This way we can reduce the clutter without fully hiding related images.

Obviously, this can produce false positives, but in our tests with various libraries we saw about one case in hundreds (if not more) of stacked images. UI to fix that could be added, although we don't think that it is sufficiently important to make the interface more complex.

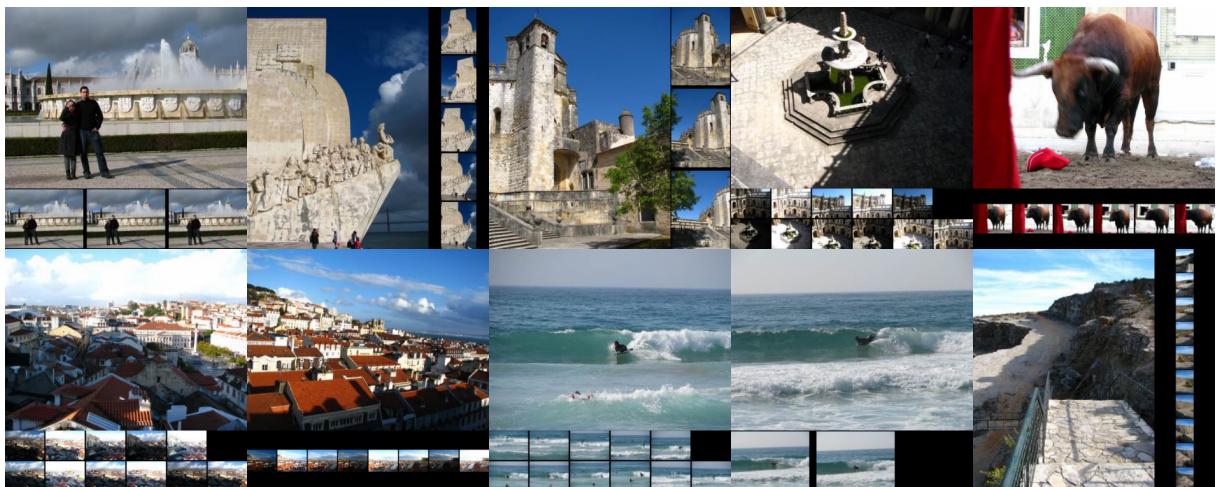


Figure 4.17: Selected examples of photo stacking as presented by Eagle Eye.

Figure 4.17 is a selection of success cases for this feature. It reduced a set of 80 images to just 10. It's possible to see cases where multiple pictures were taken in sequence and cases where the aim was to make a panorama. Portrait and landscape images are both equally well handled. The main image has the same size as it would be without the grouping, but is aligned to the top left corner, maximizing the free space to the left or to the bottom of the grid. This space is then filled with the other images which are adjusted in ways that can maximize their size.

Figure 4.18 is a real example of a trip to the beach, where some photos were taken normally, some were taken using the HDR or the Panorama techniques and some others were taken as a quick sequence. All this sets of photos are correctly grouped together. This figure contains a total of 71 photos, taking the space of only 36. On the bottom left of the figure, can be see two sets of sequential photos of the same scenes that were not stacked together due to the longer period of time between their captures. To correctly group this photos in an automatic way, their visual similarity must be determined to avoid grouping unrelated photos. This is one of the features already referred on the Solution Requirements chapter (3.3).

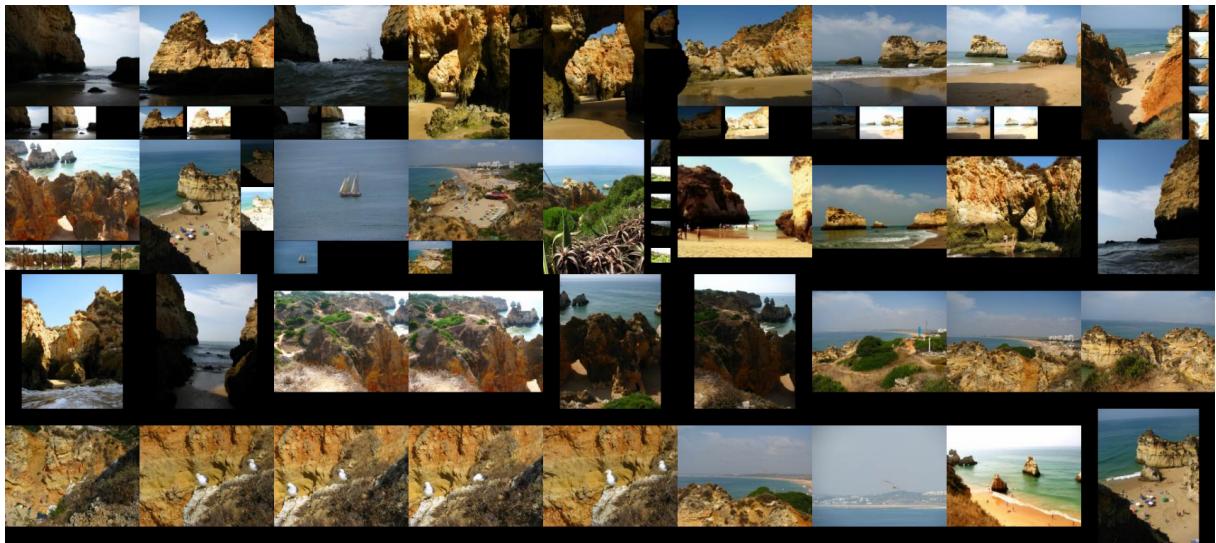


Figure 4.18: Example of a trip where HDR and Panorama techniques where applied for some photos. Other shots were repeated, some quickly and others not so much.

To conclude this chapter, where we've seen what this work does and how it does it, we would like to add that many features and improvements could be added for an even better experience, but time and the focus of this work didn't allow us to do it. Even so, the section on Future Work (6.2) contains many ideas we had which could bring this work closer to a real world application.



# **Chapter 5**

## **Evaluation**

To better understand how well Eagle Eye works, we performed user testing and, to do that, we defined a number of goals for Eagle Eye and a testing methodology, with specific tasks to be performed by the users. We also did separate tests with the users' own libraries to evaluate the benefits of knowing the contents of the system.

### **5.1 Objectives**

With our tests, we wanted to observe how various parts of Eagle Eye worked for a variety of people. We will answer the following questions, with our evaluation of the system:

- How the users react to the large quantity of images displayed simultaneously?
- How easily do the users understand how to use the navigation controls and what could be improved?
- Do users think the available disposition and sorting options are adequate and useful?
- Can users understand what are the filters capable of, or there is a need for an improvement?
- Can users extract interesting information from the collections by using Eagle Eye?

The answers will be provided later on, after explaining the testing and the results.

### **5.2 Testing Methodology**

We performed two different tests with users, both with different purposes.

One of the tests is the ideal use-case for our application: the scenario where the users are interacting with their own photos through our system. A few users gave us their photo libraries for us to process, so that we could test the usage of the user's own library in Eagle Eye.

The other test we performed allowed us to broadly analyze how users responded to the UI, how they interacted with the system, how long some tasks took to complete or how did they chose to perform a

certain task. This test only makes sense if we can replicate the conditions for each user so we generated a fixed collection, equal for all users who performed this test. This collection had 5679 photos and was based on a subset of my own photo collection, with images spanning from 2005 up to 2011. We had a larger number of users perform this test, for more accurate quantitative results.

Testing was done by presenting the user with the Eagle Eye's visualization interface ready to use. They never saw the backend or any of its processing steps, since that part is currently less than polished, takes a long time to process, and is not what we are interested in testing.

The interface and controls were explained in detail to the users, from how to navigate the collection to the functions of the buttons on the UI. They were then allowed to play around, explore and learn Eagle Eye by themselves. I was always next to them, giving little hints and some occasional help. After they felt comfortable with the system, I asked them to perform a few tasks, measured time and errors and in the end asked them to answer a few questions, describing what they thought of the system.

Most tests were performed on my personal laptop computer, connected to a 22" external monitor. The big screen is very important and helps provide a better experience while using Eagle Eye, for its capacity to display the images in a larger and more comfortable size than my 13" screen computer. Other tests had to be performed directly on the laptop, since I had to meet some users at their work locations on the university campus.

## 5.3 Characterization of users

We tested 18 people, 16 of them students in this institution, between the ages of 21 and 25, from the various degrees of computer engineering. The rest were composed from people between 40 and 55 years.

Nine of our users said they took photography seriously, although not professionally. Eight said they just like to take pictures for keeping memories or just used their camera phone and the other person doesn't take pictures (fig. 5.1.1).

Ten people claim to have less than 5000 photos stored on their computer, two people said they had between 5 000 and 10 000 while the other eight people had more than 10 000 photos (fig. 5.1.2).

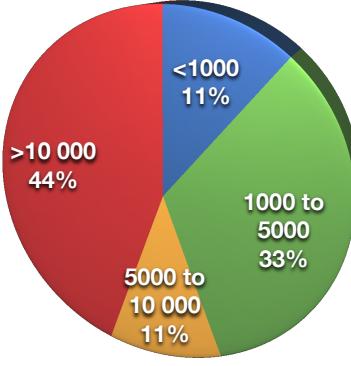
From the 18 users, only the user that doesn't take pictures relied exclusively on the organization provided by an application, while all others keep their photos on folders, some using applications to help manage them. The applications in use can be seen in fig. 5.1.3. Additionally, all users except one have a systematic way to organize their collections, being them inside an app or not.

## 5.4 Test Results

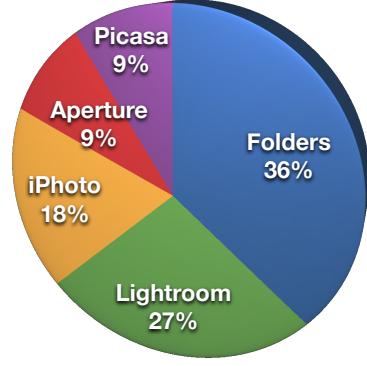
We will now go through the results of the testing, starting with the behaviors of the users during the process of learning to use Eagle Eye. After that, users were asked to perform some tasks and we will see some of the more interesting parts of that process, followed by the most common errors and



5.1.1: How users see themselves as photographers.



5.1.2: Size of users' photo library.



5.1.3: What they use for organization.

Figure 5.1: User characterization.

problems faced by the users. After the test, we asked the users to answer a inquiry and its results will be detailed towards the end of this section.

#### 5.4.1 User behavior during exploration

Upon being told how the system worked, users generally went ahead and played a while with the zooming, and viewed some images and groups at will, while noticing how the groups can be differentiated when zoomed out.

Next, users tried the different disposition options, generally in the order they are presented in the UI. The motif for two date displays was usually questioned and, after an explanation, they move on to the color display which they usually liked for the nice and clear color separation, although, when viewed more closely, some images appeared in odd locations.

After understanding the differences in dispositions and viewing some groups/images more zoomed in, the users tried the filtering or the manual selections for reducing the amount of images displayed and focus only on part of the collection. The filter showed useful even though its bugs occasionally made an appearance. The manual selections were easier to use and for that captivated more usage from the users.

#### 5.4.2 Results for tasks

After playing around with Eagle Eye, we asked the users to perform some tasks for us. We measured how long did they take to do each task and what errors they've done. Below the description of the tasks are graphs (figs. 5.2 and 5.3) with a summary of the measurements.

We started with simple tasks, to see if they understood the basics of the system. Tasks like "Switch to the color view", "Select and display a group", "Select and display a few images" weren't problematic and generally took between a couple seconds up to about 15 seconds, depending on the user, and without errors. We also asked for a simple combination of functions with the task "what is the most common color of the pictures taken with the most used camera?", which is the process of sorting by camera,

selecting a group and sorting by color. 18 people finished the test in less than half a minute without problems, the other two made one mistake

Next, we asked users to select images from the year 2010. This could be done by using the filter to only display images of that year or by viewing images by sequential date and manually select the images that corresponded to that year. Most users tried typing 2010 in the filter bar and quickly got to the intended result in about 5 seconds, but four users still looked into the default, treemap-based, date view until they remembered/noticed it wasn't sorted and then looked for the sequential date view and performed a manual selection of the images. These users took about a minute to perform the operation.

In another test, we told the users that the library had photos of a certain annual event and asked them to figure out what years had photos of that event. The easier way was to filter by the event name, sort by date and view the group names, which would take less than 15 seconds. Slightly more than half the users did this method. Others went to the path view and looked for groups containing the event name, taking between one and two minutes to perform the task. A small percentage tried to discover the photos by navigating on the complete image set, taking an even longer time to complete.

The users were then asked to find a photo of a jumping dolphin. This test could be accomplished in variety of ways. Some users had already noticed a group of photos of the Zoo and zoomed into those, quickly finding the photos in less than ten seconds. Others changed the view to be grouped by colors and selected only the blue groups. The filter returned the photos in no apparent order and, after complaining about that, the users changed the view to date or path, finding the group of Zoo photos on the corner. Some zoomed in, others selected that group and found the photos, not taking less than 30 seconds to do it. A few users tried at first to see if there were some dolphin-related keywords but, unfortunately, those photos had no keywords associated.

### 5.4.3 Errors and misunderstandings

During the above two tests, exploration and tasks, some errors occurred somewhat repeatedly, some caused by bugs, some by poor choices in the UI.

Users could not quickly understand the difference between the buttons “Date” and “Date (Linear)” buttons by themselves making some mistakes sometimes.

A few users thought the “Select images” button was for applying the filters and not for manually selecting images, generating some confusion.

A few other users complained about the way to manually select a single image. Currently, the user is required to make a “click-and-drag” selection to select one or multiple images, conflicting with the common way of “click to select item; click-and-drag to select many”.

Another problem, detected by a few users, that is also related to selections is the action of de-selecting images and groups, which can only be performed by deleting the entry in the filter bar, and not by interacting with the images.

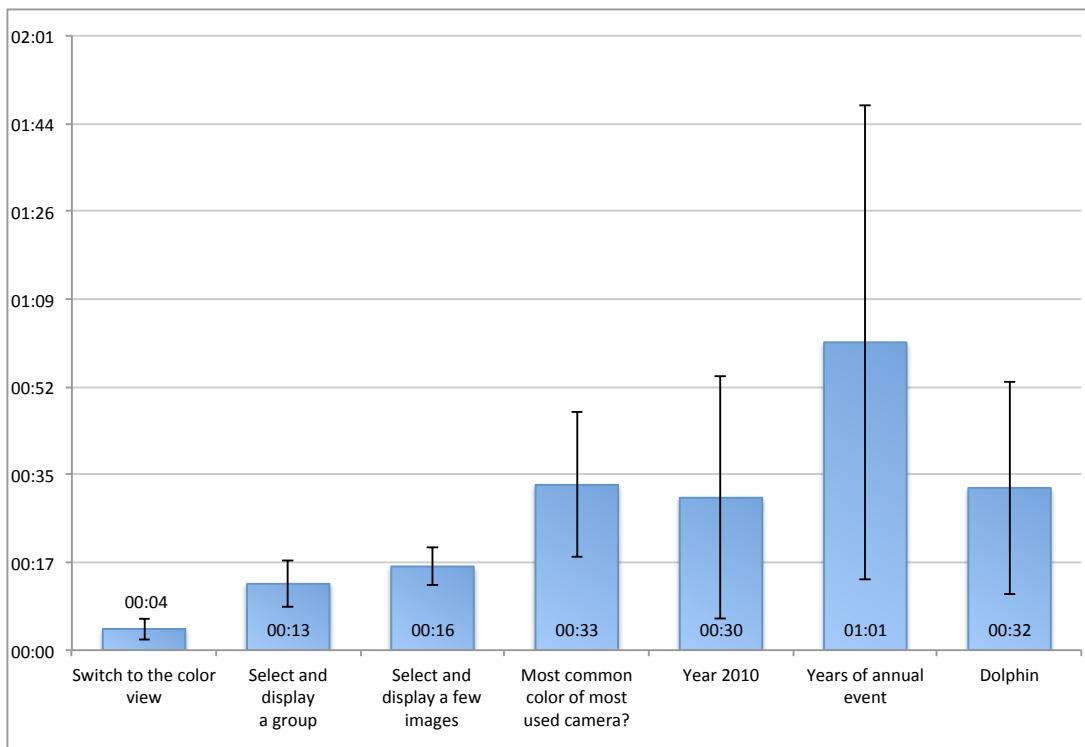


Figure 5.2: Time (average and standard deviation) spent by users for each of the tasks.

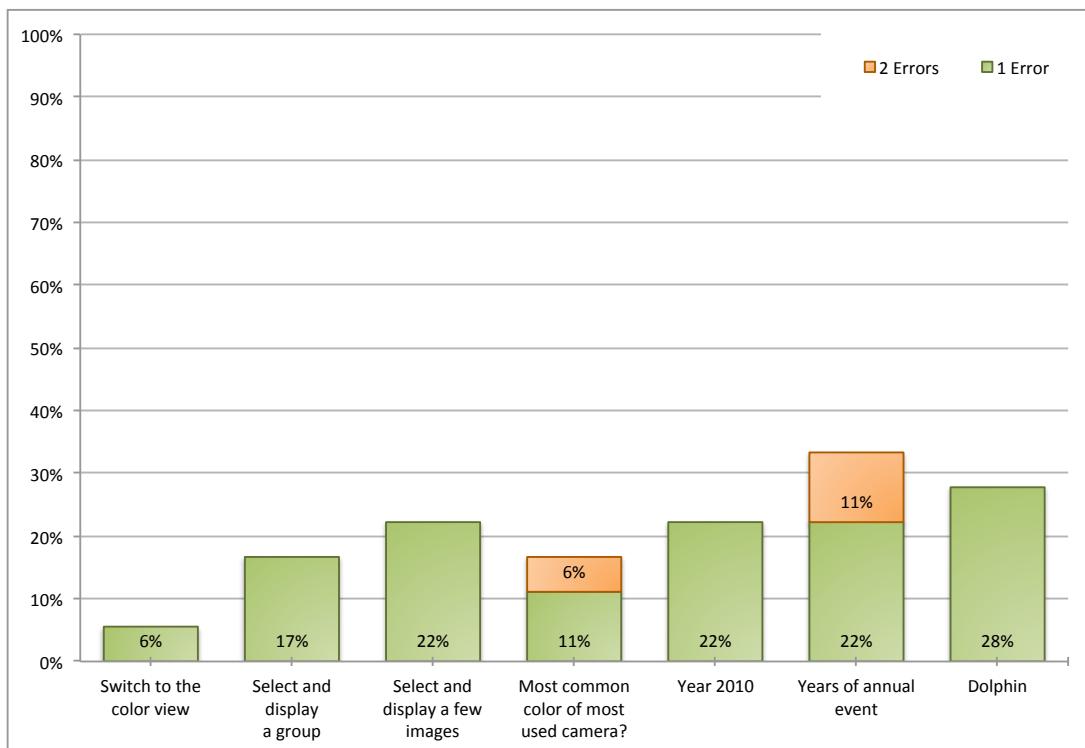


Figure 5.3: Percentage of users who made errors on each task, grouped by how many errors.

#### 5.4.4 System Usability Scale

On the inquiry filled by our test users after the testing the system, we had a section with the System Usability Scale (SUS) questions. SUS[5] is a simple and generic ten question scale, giving a global view of the usability of a system. The SUS result of inquiring 18 people was 75 points<sup>1</sup> on a scale of 0-100 points, which is a fairly good mark for our system. With the missing improvements it could score higher.

The averaged results to each of the answers of SUS can be seen on Table 5.1. The questions' focus alternates between positive and negative characteristics for the system. Some important questions got good rates, like the “easy to use” and “quick to learn” with ratings of 4.3 and 3.9 of a maximum 5 and the “unnecessarily complex” question also getting the second best rate for a negative question, 1.6, which means users generally found the system to be easy to use and learn, although it still has some margin for improvement on the user experience area. Another important question is the “I would like to use this system frequently” which got a rating of 3.4, slightly lower from what we would like to see, yet still above the waterline. We tried to understand if the kind of photography or the size of the library had an influence on this value, but we got mixed results and concluded that they probably aren't connected. Even though, this rating reinforces the point above about the need for improving the system, but this time by adding features that could make the user want to use the system more often.

I think that I would like to <b>use this system frequently</b>	3.4
I found the system <b>unnecessarily complex</b>	1.6
I thought the system was <b>easy to use</b>	4.3
I think that I would <b>need the support of a technical person</b> to be able to use this system	1.6
I found the various functions in this system were <b>well integrated</b>	3.5
I thought there was <b>too much inconsistency</b> in this system	1.9
I would imagine that most people would <b>learn to use this system very quickly</b>	3.9
I found the system very <b>cumbersome</b> to use	1.9
I felt very <b>confident</b> using the system	3.7
I needed to <b>learn a lot of things</b> before I could get going with this system	1.5

Table 5.1: Average of the results for each of the ten questions of SUS. Values range between 1.0 (strongly disagree with the question) and 5.0 (strongly agree). Gray rows are “positive” questions, where higher values are better, and white rows are “negative” questions, where lower values are better.

#### 5.4.5 What users liked

Questioned of what they liked on the system, and from what was seen from the interaction, we can confidently say users liked the general concepts of the system. I now present a summary of the users opinions of the good points of Eagle Eye:

<sup>1</sup>With a standard deviation of 9 and a median of 75.

- It was easy and simple and it helps to get a good overview of the photo collection;
- the photo display of a much greater amount of photos than what they've ever seen, making visualization of many images easier;
- the easy zoom in and out to quickly see groups or photos in full size or in smaller size;
- still being able to identify groups of images which are small sized;
- the intersection of filters/groups, i.e., the various ways to quickly organize and access the images, specially the ability to view groups by events or colors;
- the ability to select images;
- the smart use of space by grouping related images;

Essentially, people seem to like the new interaction methods to browse photos, specially if they know the photos themselves, which makes the overall usage of the system much more interesting.

#### **5.4.6 User suggested improvements**

During the course of the tests and on the questionnaire, some users suggested some improvements to our work. Some of these improvements were already in our mind, but we didn't have enough time to implement them. We will now expose the suggestions we agree on:

- Information to help the user use the system, e.g., help, guides or tutorial;
- When there's many images, display a few from each group in a bigger size and on zoom in, display more;
- Easier way to quickly view a single image/group in full size, like a double click;
- The filter should be improved, removing the bugs, giving some guidance to the user and allowing both AND and OR logics (currently only uses the former) and a button to clear it;
- The selection buttons should be separated from the filter bar to avoid misleading people;
- The display should be organized hierarchically where appropriate, e.g., date, path and color based visualizations;
- The color measure should be improved;
- Event-based visualization, better than just dates or paths;
- Face detection should be improved to Person recognition;
- Improved method to select single images;
- Display the number of photos on screen;
- Better group separation.

Although we tried to make the user interface as clean and simple as possible, a couple users found some difficulties and suggested some simple guides. We think a simple introduction to the system is nice and helps the users, but it wouldn't make sense to spend time doing it at this point.

## 5.5 Accomplished Objectives

We can now answer what objectives for Eagle Eye were accomplished.

### **How the users react to the large quantity of images displayed simultaneously?**

Users generally liked the large view and were surprised by how much can still be recognized, specially when working with their own photos and if they're grouped by events.

### **How easily do the users understand how to use the navigation controls and what could be improved?**

Users like the navigation controls, however many felt it should be easier to go directly to a full view of one image and some agreed that keyboard navigation or slideshow-like controls should make the navigation easier.

### **Do users think the available disposition and sorting options are adequate and useful?**

Users liked the ability to switch between dispositions and preferred some more than others. Users generally liked the date display in groups, but some would prefer that it would be more faithful to the time. Some suggested that it would be better if the tree map was organized with groups and subgroups so that all photos of a year/month got grouped together. Users understood the current need for the linear date display, yet found it hard to view anything and didn't spend much time with it. The other preferred disposition was the color, the overall result of the grouping was pleasing to the eye of most users. The problem was when they started zooming in and some images revealed to be a mix of different colors from the color of the group they were in. This was a known problem of the method used. Another problem that revealed was the large size of groups like "Gray" or "Black" usually took great part of the screen, leaving all the more interesting colors to half of the screen. Some users even filtered out this black and grey groups by selecting and viewing only the colorful ones. The path display wasn't very used since it ended up being much similar to the date display, except on some libraries where folders weren't named by date, but by event and made some sense to view the groups by these names.

### **Can users understand what are the filters capable of, or there is a need for an improvement?**

Some users understand what the filter is capable of when the autocomplete gives some suggestions, others could use some initial guidance to understand that a lot of things can be used on it. The biggest problem is that some of the information they would like to use can't be obtained automatically, like location (unless already on EXIF) or names of people, unless we hook into some third party app where the user have previously saved this data.

## Can users extract interesting information from the collections by using Eagle Eye?

Users can extract some information about their collections, like collection sizes, days with many shots, most common colors, most common devices and such. It would be even more useful with the inclusion of a second grouping and sorting option so, for instance, a disposition by devices could have the images inside each camera sorted by their color. This way even more information could be retrieved from the collection.

## 5.6 Case Studies

As referred on the beginning of the chapter, we made some tests with some users' photo collections. We managed to get five users willing to give us a copy of their collections for a more personalized testing.

We got libraries containing from 833 up to 6064 images. None of the tested users report being able to recognize most of the groups in Date view. The user with the most photos said he needed to zoom in a little to understand better some of the groups, since many of them had a similar, grayish, tone that made them a bit hard to distinguish. We experimented filtering out all the images in the gray category using the color view and get back to the date view and groups were much easier to distinguish, because there were a lot of less images to view, making the visible ones bigger, and because groups had now a more relevant color. Although we were hiding images, the view itself ended up more interesting.

Most tested users only had a couple images stacked and did not use any of the techniques that produced sequential images (e.g. panoramas, bracketing, burst), but the person with the biggest library did make some bursts and some long time lapses and most were correctly stacked. This stacking allowed his collection of over 6000 photos to fit on the space occupied by a collection of only 4368 images, which is a 27% reduction of clutter, without actually hiding the images. As an example, specific to this library, we selected the largest stacks on his library and compared how much space they took compacted and expanded. The six stacks we chose (fig. 5.4) contained, in total, 563 images. By stacking, the 563 very identical images took only the space of 6, which is about 1% of the space. We know that not many people make this type of photography but, for those who do, stacking is a very helpful feature which can only be improved.



Figure 5.4: Case study: six time lapses of 563 images, reduced to 6 spaces on the canvas

## 5.7 Discussion

With the tests we performed, we verified people generally liked Eagle Eye and saw the potential of the system.

We designed it to be simple while powerful, and people proved it was simple, by their positive answers and by their usage of the system, where we not only measured generally low times to complete the tasks, but also a low number of errors. We acknowledged some aspects of the interface could be improved for a better and more user-friendly interaction.

By the user interactions and suggestions, we also understood that Eagle Eye could be even more powerful, and with some improvements on the extraction of features and by enhancing the capabilities of the organization on canvas, it could be much more interesting and relevant to anyone who keeps photos on their computer.

# **Chapter 6**

## **Conclusions**

We will now conclude this work and go through what we think could be improved.

### **6.1 Conclusion**

Photo visualization is an eternal problem with lots of compromises. Big software companies sell the simple ways with big lists of photos, but many other ways have been explored for years, some more eccentric than others. Some more useful than others. We explored their good and bad points and selected the features we thought important. Made a user survey and fine tuned our vision and Eagle Eye was born.

Our work aimed at being useful to users who have a regular collection of photos on their computers and want a better way to view and learn more about them. We built a system that gives users a completely different view of their collection, by showing everything at the same time. While overwhelming at first, we have seen that users like it and, by already having knowledge about their own library, they can understand each group of photos, whether their photos were taken yesterday or ten years ago. From this, we allow users to change, reorder, filter, zoom and enjoy their pictures.

For this to be possible, we developed a system that extracts information from images and that can be expanded with more processes to obtain interesting and uncommonly seen information. This allows users to observe their collection in a variety of ways, by selecting different information and filtering down to specific sets of images.

Although we didn't had time to implement many features we had in mind for an even better system, our user tests revealed that people liked the system, giving it a moderately good rating of 75 in a scale of 0 to 100.

We think we have a solid concept that, with some more development, could be made as an application that could actually be released to the public.

## 6.2 Future Work

During the realization of this work, many ideas popped in our minds but we didn't had the time to work on them.

### 6.2.1 Consolidation

The most important effort that must be done to this work is a rearrangement of the backend and its connection to the visualization.

As we've seen on chapter 4.2, both parts of this work are currently separated, with the backend being a command line utility. Our idea was to create a graphical application where the user could enter the folder paths he wants to add to Eagle Eye and then the application would do the needed work on the background, using free resources, and when it was done, the user could enter in the visualization mode without the need to open the internet browser. Better yet if the system allowed things to be viewed while the processing was still running.

This interface would also provide some configuration options, for the system and its plugins.

### 6.2.2 Feature Extraction Plugins

The second important thing to do is to improve the feature extraction plugins.

The color plugin, for instance, should be able to better identify colors in images instead of simply make a color average.

More data should also be extracted from the EXIF. As example, location data should be translated into place names<sup>1</sup>. Another plugin we see some benefit in creating would be a generator of keywords related to the capture dates. Currently we have the visualization reading the capture dates and generate keywords like "Summer" or "May 2011" for use as search tags. By transferring this to a feature extraction plugin, this data will be generated only once. It would then appear in the Suggestions of the Filter Bar as date options, allowing the gathering of sets of images from various years or seasons.

An additional plugin that would be interesting to develop would be some kind of hook into other application's metadata about photos. As an example, it's possible to access Picasa's face recognition and identification data about photos. By bringing this data into Eagle Eye, we could allow the user to perform searches for people's names without having to add those names as keywords.

### 6.2.3 Visualization

On the visualization part, the linear disposition, used mainly by the date view, should be improved to make more clear the distinction between days, months, years. A great idea that would be interesting to use would be the inclusion of something like the date bar in the work of Grgensohn et al. [27] (fig. 2.10).

---

<sup>1</sup>The action of turning GPS coordinates into addresses is called reverse geocoding. An example of a provider of such service: <http://code.google.com/apis/maps/documentation/geocoding/index.html#ReverseGeocoding>

Another idea for a different disposition born from the stacking of similar photos discussed in section 4.3.3, where one image is bigger than the rest, to help identify what is on what group.

We also think it needs to be more slideshow friendly. Slideshows are important part of any photo application, be it on the desktop or even on the web, since people always like to view their images. This work already displays images in fullscreen, it just needs a better interaction to repeatedly move to the next image, something that can be achieved easily with a button on the UI and bindings to keyboard keys.

On our evaluation, users suggested a number of things they would like to see (in 5.4.6) and, after consideration, some could very well be included.

#### **6.2.4 Connect**

Another interesting point that would make the system stand out would be the connection between services like calendaring or location-tracking. Crossing their informations with our library could allow searching for photos of a certain event or location without the need for the library to contain that specific information. For instance, we could search for “Mary’s wedding” on our application and, by looking into the calendar, find out what that day was and display the related photos.



# Appendix A

## User Survey

We asked some people to answer a survey and we got 30 answers that helped us understand photography related habits and user wishes, and compiled them in this appendix.

### A.1 User characterization

Our survey started by asking the respondents to describe themselves and their libraries.

#### Kind of Photographer

We asked the respondents what of the following options described them the best:

- I'm a **professional** photographer.
- I like to take photography **seriously** but I'm no professional.
- I like to take pictures just to keep as **memories**, but I'm not a photo-geek.
- I only use my camera **phone** for fun.
- I **don't** take pictures.

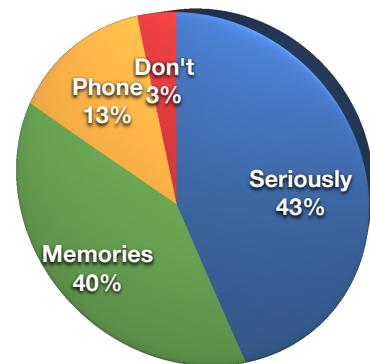


Figure A.1: Kinds of photographers.

As we can see on fig. A.3, most respondents like photography and take photos, being mostly people who take photos for memories and people who are really into photography but are not professionals.

#### Storage of digital photos

The second question was a check to verify if the respondents stored digital photos in their computers. The questionnaire would end if they said no. Fortunately, everyone said yes.

## A.2 Characterization of Photo Library

We then inquired about their photo libraries.

### Library size

To start the survey about the library, we inquired how large was the respondents photo library.

We made this an open question, because we wanted to obtain real values from the users and not just limiting to a number of options that hardly represent anything. The answers for this question can be seen on fig. A.3.

We were surprised by the values we got, going from 5 up to a million photos.

Almost half of the answers are between the 1 000 and the 10 000 values, having two thirds of the results falling between 0 and 10 000. We also found the median value of the answers, which is 5 000.

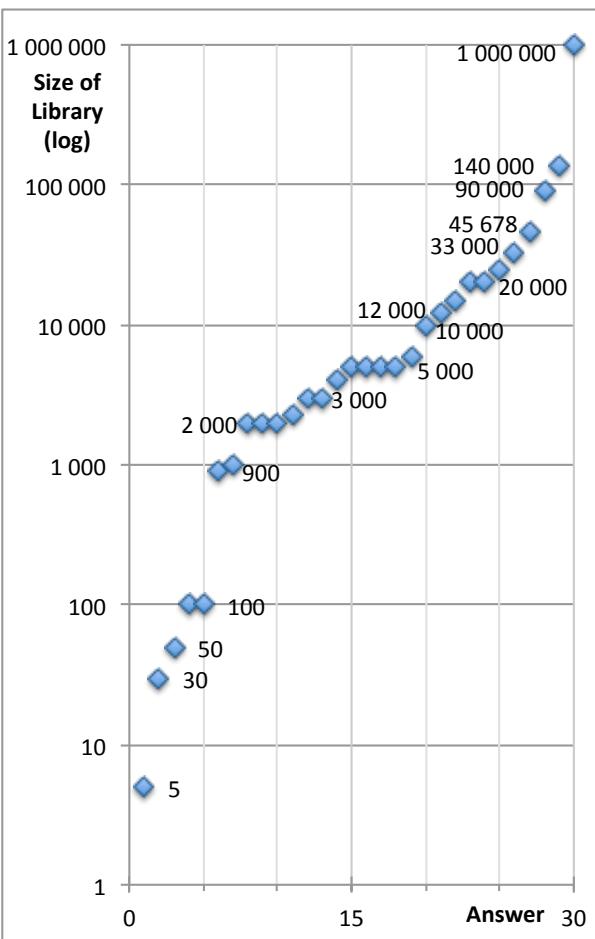


Figure A.2: Library Size

### Usage of Software

We then questioned about what software applications did people use to help them organize their collections. Almost half admitted they only used folders for their organization, while the other 17 people used some kind of software.

Going through each software applications, we can see that 40% use an entry-level software (iPhoto and Picasa) while 17% prefer more Pro-level systems.

We tried to correlate the size of the collections with the software used but we realized that we probably don't have enough data for some of those softwares and that users might not have their entire collection inside the application, therefore defeating the correlation.

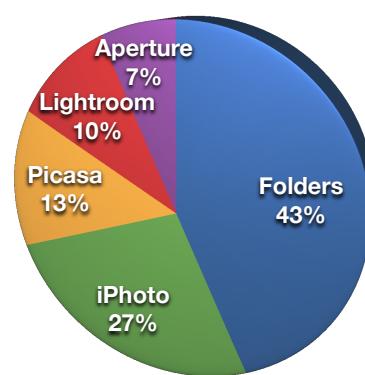


Figure A.3: Software used

## Systematic Organization

The question “Do you have a systematic way to keep your photos organized?” is meant to understand if people keep a system for their library, either on folders or inside an application.

Since most of the applications in the last question rely on folder organization, is no surprise that having a system in place for keeping things organized in folders is the clear winner of this question. Five people answered that they have this system inside an application and correlating this question with the previous one we found another expected aspect: only entry-level software is used to keep everything tidy, without touching folders, more exactly four people say they use iPhoto and one says Picasa is his/her choice.

In contrast, 20% of the respondents claimed not having a way to organize their collections. The median of the number of photos of this group of people is 1050, the lowest, when compared to the other groups, the ones who use applications and the ones who rely on folders, both with 5000 of median.

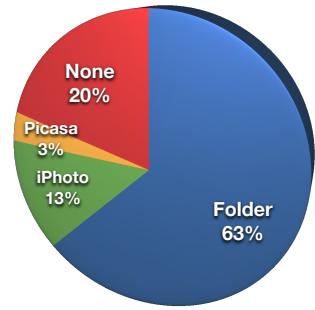


Figure A.4: Systematic organization in folders, in an application or no organization.

## Photos in Close Sequence

The following question was intended to see how many users take photos in close sequence, e.g., bursts, panoramas, bracketing or time lapses, where the photographer ends up with many photos in a few seconds.

As we can see on fig. A.5, 40% of the respondents said they do take photos in close sequence. We also saw that this group has a median value of library size of 8000, whereas the other group stays on the 2150 images. In fact, of the total 13 enthusiasts, 11 of them answered yes on this question, followed by someone who claims only to take photos for memories.

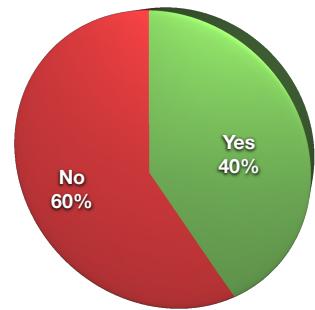


Figure A.5: Users that take photos in close sequence.

## Preferred Features

We had some ideas of what we would like to see extracted from the images and exposed to the user via our application and we asked for people’s opinion on this survey and the results can be seen on fig. A.6.

While we weren’t surprised by the favorites Date/Time, Event and Location, we did expect some higher ratings for number of faces and colors. Orientation and camera settings got the ratings we did expect.

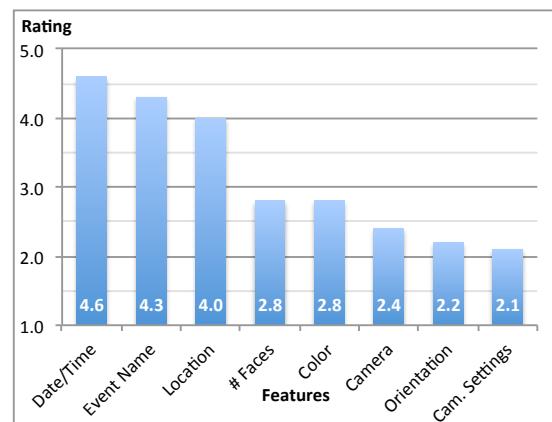


Figure A.6: User rated features.

## Suggestions of Features

In line with the previous question, we asked what other features they would like to see extracted from images.

- The most common request, with about four different people requesting it, is people recognition. This allows photos to be tagged with the names of people in the photo and not just how many faces there are on the photo.
- Someone asked that, from a set of similar photos (exemplified with friend group photos), pick the one that, e.g., is more clear, has a better color, more smiles, more open eyes.
- Metadata was requested by three people, e.g., EXIF metadata, tags, format (RAW/JPEG), if the photo is the original or the edited version, the number of duplicates, the location on disk.
- One person asked for identification of whether the photo was taken during the day or the night.
- Another person was interested in having information about objects, buildings or animals.

## General Overview

To conclude the survey, we asked if the respondents feel they have the tools to have a good overview of their collection at the same time.

The answers to this question rose some uncertainty, due to their mixed views. Each person is different from the next and so are their requirements. For instance, some people considered iPhoto/Aperture's events or projects as good overview for the library, but others don't think the same way about it, claiming it is limiting. On fig. A.7, is a summary of whether people think they can get a good overview or not, split by the applications (or lack of them) they use. Also included is a global percentage for comparison.

One thing we can be sure, though: those who use folders don't think they can get a good overview from them.

With this we conclude our survey.

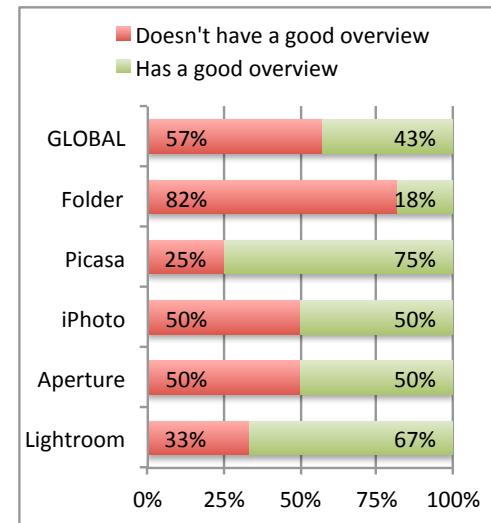


Figure A.7: Whether people think they can get a good overview or not.

# Bibliography

- [1] ALBUZ, E., KOCALAR, E., AND KHOKHAR, A. A. Scalable Color Image Indexing And Retrieval Using Vector Wavelets.
- [2] AZZAM, I., LEUNG, C., AND HORWOOD, J. A fuzzy expert system for concept-based image indexing and retrieval.
- [3] BEDERSON, B. B. PhotoMesa: a zoomable image browser using quantum treemaps and bubblemaps. In *Proceedings of the 14th annual ACM symposium on User interface software and technology* (2001).
- [4] BERMAN, A., AND SHAPIRO, L. A Flexible Image Database System for Content-Based Retrieval. *Computer Vision and Image Understanding* (1999).
- [5] BROOKE, J. SUS: A quick and dirty usability scale. In *Usability evaluation in industry*. 1996.
- [6] BRUIJN, O., AND SPENCE, R. Rapid serial visual presentation: a space-time trade-off in information presentation. *AVI '00: Proceedings of the working conference on Advanced visual interfaces* (2000).
- [7] BRULS, M., HUIZING, K., AND WIJK, J. J. V. Squarified Treemaps.
- [8] CAMPBELL, I., AND VAN RIJSBERGEN, C. The ostensive model of developing information-needs. *Citeseer* (2000).
- [9] CHANG, S., YAN, C., AND DIMITROFF, D. An intelligent image database system. *IEEE Transactions on ...* (1988).
- [10] CHANG, Y.-C., AND LEE, D.-J. Color Image Quantization Using Color Variation Measure. In *Proceedings of the 2007 IEEE Symposium on Computational Intelligence in Image and Signal Processing* (2007).
- [11] CHEN, J.-Y., AND BOUMAN, C. A. Hierarchical browsing and search of large image databases. *Image Processing* (2002).
- [12] CHEN, J.-Y., BOUMAN, C. A., AND DALTON, J. C. Similarity pyramids for browsing and organization of large image databases. *Human Vision and Electronic Imaging III* (1998).

- [13] CHEN, L., HU, B., ZHANG, L., AND LI, M. Face annotation for family photo album management. *International Journal of Image ...* (2003).
- [14] CHRISTMANN, O., AND CARBONELL, N. Browsing through 3D representations of unstructured picture collections: an empirical study. *Proceedings of the working conference on Advanced visual interfaces* (2006).
- [15] COMBS, T. Does zooming improve image browsing? ... of the fourth ACM conference on ... (1999).
- [16] COOPER, K., BRUIJN, O., SPENCE, R., AND WITKOWSKI, M. A comparison of static and moving presentation modes for image collections. *AVI '06: Proceedings of the working conference on Advanced visual interfaces* (2006).
- [17] COOPER, M., FOOTE, J., AND GIRGENSOHN, A. Temporal event clustering for digital photo collections. *ACM Trans. Multimedia Comput. Commun. Appl.* (2005).
- [18] CUNNINGHAM, S., AND MASOODIAN, M. Identifying personal photo digital library features.
- [19] DANIELS, J., HA, L., AND OCHOTTA, T. Robust smooth feature extraction from point clouds. *computer.org* (2007).
- [20] DATTA, R., JOSHI, D., LI, J., AND WANG, J. Image retrieval: Ideas, influences, and trends of the new age. *ACM Computing Surveys (CSUR)* (2008).
- [21] DATTA, R., LI, J., AND WANG, J. Z. Content-Based Image Retrieval - Approaches and Trends of the New Age. In *MIR '05 Proceedings of the 7th ACM SIGMM international workshop on Multimedia information retrieval* (2005).
- [22] DENG, J., DONG, W., SOCHER, R., LI, K., AND FEI-FEI, L. Imagenet: A large-scale hierarchical image database. *computer.org* (2009).
- [23] FLEMING, M. Categorization of faces using unsupervised feature extraction. *Neural Networks* (1990).
- [24] FOGARTY, J., TAN, D., KAPOOR, A., AND WINDER, S. CueFlik: interactive concept learning in image search. *Proceeding of the twenty-sixth annual SIGCHI conference on Human factors in computing systems* (2008).
- [25] FONSECA, M., AND JORGE, J. NB-Tree: An indexing structure for content-based retrieval in large databases. *Proc. of the 8th Int'l Conf. on Database Systems for Advanced Applications. Kyoto: IEEE Computer Society* (2003).
- [26] GABBOUJ, M., BIRINCI, M., AND KIRANYAZ, S. Perceptual color descriptor based on a spatial distribution model: Proximity histograms. In *2009 International Conference on Multimedia Computing and Systems (ICMCS)* (2009).

- [27] GIRGENSOHN, A., SHIPMAN, F., TURNER, T., AND WILCOX, L. Flexible access to photo libraries via time place, tags, and visual features. In *Proceedings of the 10th annual joint conference on Digital libraries* (2010).
- [28] GIRGENSOHN, A., SHIPMAN, F., WILCOX, L., TURNER, T., AND COOPER, M. MediaGLOW: organizing photos in a graph-based workspace.
- [29] HEESCH, D. A survey of browsing models for content based image retrieval. *Multimedia Tools and Applications* (2008).
- [30] HEESCH, D., AND RÜGER, S. NNk networks for content-based image retrieval. *Advances in Information Retrieval* (2004).
- [31] HILLIGES, O., BAUR, D., AND BUTZ, A. Photohelix: Browsing, Sorting and Sharing Digital Photo Collections. *Horizontal Interactive Human-Computer Systems, International Workshop on* (2007).
- [32] HONG, Z. Algebraic feature extraction of image for recognition. *Pattern Recognition* (1991).
- [33] HSU, R.-L., ABDEL-MOTTALEB, M., AND JAIN, A. K. Face detection in color images. *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2002).
- [34] Hsu, S., CUBAUD, P., AND JUMPERTZ, S. Phorigami: A Photo browser based on meta-categorization and origami visualization. *Human-Computer Interaction. Novel Interaction Methods and Techniques* (2009).
- [35] JAIN, A., AND VAILAYA, A. Image retrieval using color and shape. *Pattern Recognition* (1996).
- [36] JOHNSON, B., AND SHNEIDERMAN, B. Treemaps: a space-filling approach to the visualization of hierarchical information structures. In *Proceedings of the 2nd conference on Visualization '91* (1991).
- [37] KRISHNAMACHARI, S., AND ABDEL-MOTTALEB, M. Image browsing using hierarchical clustering. *Computers and Communications, 1999. Proceedings. IEEE International Symposium on* (1999).
- [38] LEE, D.-J., MITRA, S., AND KRILE, T. F. Analysis of sequential complex images, using feature extraction and two-dimensional cepstrum techniques. *Journal of the Optical Society of America A* (1989).
- [39] LIU, Y., ZHANG, D., AND LU, G. A survey of content-based image retrieval with high-level semantics. *Pattern Recognition* (2007).
- [40] MINKA, T., AND PICARD, R. An image database browser that learns from user interaction. *Master's thesis* (1996).
- [41] MINKA, T., AND PICARD, R. Interactive learning with a “Society of Models”. *Computer Vision and Pattern Recognition, 1996. Proceedings CVPR '96, 1996 IEEE Computer Society Conference on* (1996).

- [42] MOJSILOVIC, A., HU, H., AND SOLJANIN, E. Extraction of perceptually important colors and similarity measurement for image matching, retrieval and analysis. *IEEE Transactions on Image Processing* (2002).
- [43] NAAMAN, M., SONG, Y., AND PAEPCKE, A. Automatic organization for digital photographs with geographic coordinates. *Digital Libraries* (2004).
- [44] NIXON, M. Feature extraction and image processing. *books.google.com* (2008).
- [45] PANCHANATHAN, S., PARK, Y. C., KIM, K. S., KIM, P. K., AND GOLSHANI, F. The role of color in content-based image retrieval. In *Proceedings. 2000 International Conference on Image Processing* (2000).
- [46] PLATT, J., AND CZERWINSKI, M. Phototoc: Automatic clustering for browsing personal photographs. ... *Communications and Signal* ... (2004).
- [47] PORTA, M. Browsing large collections of images through unconventional visualization techniques. *AVI '06: Proceedings of the working conference on Advanced visual interfaces* (2006).
- [48] QIU, G., MORRIS, J., AND FAN, X. Visual guided navigation for image retrieval. *Pattern Recognition* (2007).
- [49] QIU, G., YE, L., AND FENG, X. Fast image indexing and visual guided browsing. ... *Workshop on Content-Based Multimedia Indexing*.
- [50] RODDEN, K. How do people organise their photographs. *BCS IRSG 21st Ann. Colloq. on Info. Retrieval* ... (1999).
- [51] RODDEN, K., BASALAJ, W., SINCLAIR, D., AND WOOD, K. R. Does organisation by similarity assist image browsing? *Proceedings of the SIGCHI conference on Human factors in computing systems* (2001).
- [52] RODDEN, K., AND WOOD, K. R. How do people manage their digital photographs? *Proceedings of the SIGCHI conference on* ... (2003).
- [53] RUI, Y., HUANG, T., AND CHANG, S. Image Retrieval: Current Techniques, Promising Directions, and Open Issues\* 1. *Journal of visual communication and image* ... (1999).
- [54] SCHAEFER, G. A next generation browsing environment for large image repositories. *Multimedia Tools and Applications* (2010).
- [55] STRONG, G., AND GONG, M. Organizing and browsing photos using different feature vectors and their evaluations. *CIVR '09: Proceeding of the ACM International Conference on Image and Video Retrieval* (2009).
- [56] SURAL, S., QIAN, G., AND PRAMANIK, S. Segmentation and histogram generation using the HSV color space for image retrieval. In *Image Processing. 2002. Proceedings. 2002 International Conference on* (2002).

- [57] TAMURA, H., AND YOKOYA, N. Image database systems: A survey. *Pattern Recognition* (2002).
- [58] TAN, K., OOI, B., AND YEE, C. An evaluation of color-spatial retrieval techniques for large image databases. *Multimedia Tools and Applications* (2001).
- [59] TERROSO, E. M. *An Image Organizer with Content-Based Image Retrieval*. PhD thesis, Universidade Federal do Rio Grande do Sul, 2008.
- [60] TORRALBA, A., FERGUS, R., AND FREEMAN, W. T. 80 Million Tiny Images: A Large Data Set for Nonparametric Object and Scene Recognition. *IEEE Trans. Pattern Anal. Mach. Intell.* (2008).
- [61] VASCONCELOS, N., AND LIPPMAN, A. A multiresolution manifold distance for invariant image similarity. *IEEE Transactions on Multimedia* (2005).
- [62] WAN, S., JIN, P., AND YUE, L. An Effective Image Retrieval Technique Based on Color Perception. In *Sixth International Conference on Image and Graphics (ICIG)* (2011).
- [63] WOJCIECH, K., BASALAJ, W., AND SINCLAIR, D. Evaluating a Visualisation of Image Similarity as a Tool for Image Browsing. *the IEEE Symposium on ...* (1999).
- [64] WORRING, M. Interactive access to large image collections using similarity-based visualization. *Journal of Visual Languages & Computing* (2008).
- [65] YUNG, M. A. *Content-Based Image Retrieval*. PhD thesis, Clayton School of Information Technology, 2006.
- [66] ZHANG, R., AND ZHANG, Z. M. Addressing CBIR Efficiency, Effectiveness, and Retrieval Subjectivity Simultaneously. In *MIR '03: Proceedings of the 5th ACM SIGMM international workshop on Multimedia information retrieval* (2003).

